

# Advanced Operating Systems (6G7Z1004)

## Lecture: Concurrency

Soufiene Djahel

Office: John Dalton E114

Email: [s.djahel@mmu.ac.uk](mailto:s.djahel@mmu.ac.uk)

Telephone: 0161 247 1522

Office hours: Monday 13 -15, Tuesday 11-12

# Recap from last week

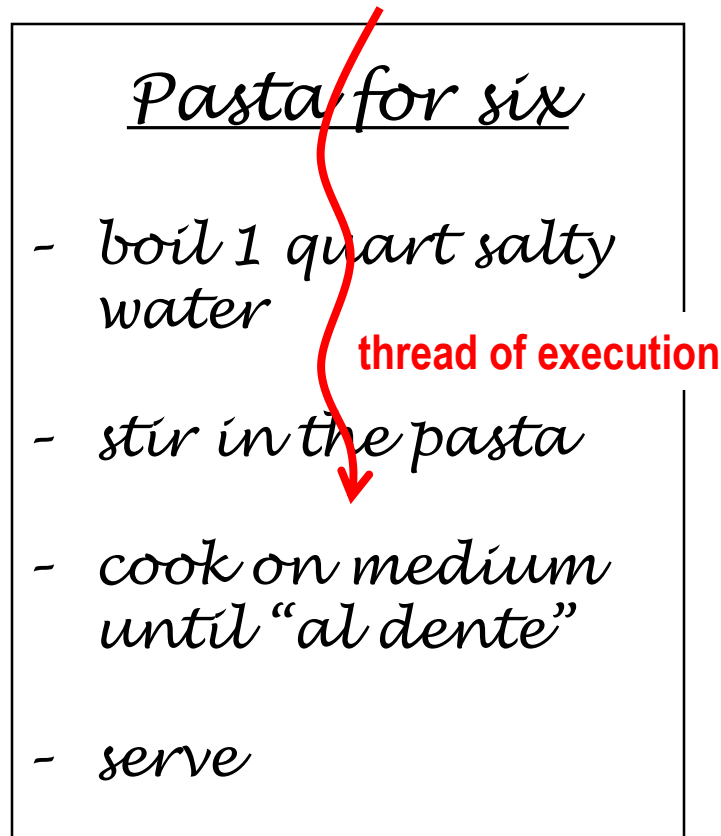
- OS principles
- Processes, threads and scheduling
- Concurrency
- Memory management
- File system management
- Privacy and Security

- a. Processes (Summary)
  - ✓ What is a process?
  - ✓ Process states
  - ✓ Process description
  - ✓ Process control
- b. Threads (Summary)
- c. Concurrency
- d. Deadlock

# a. Processes (Summary)

## b. What is a process?

- A process is the activity of executing a program



Program



CPU

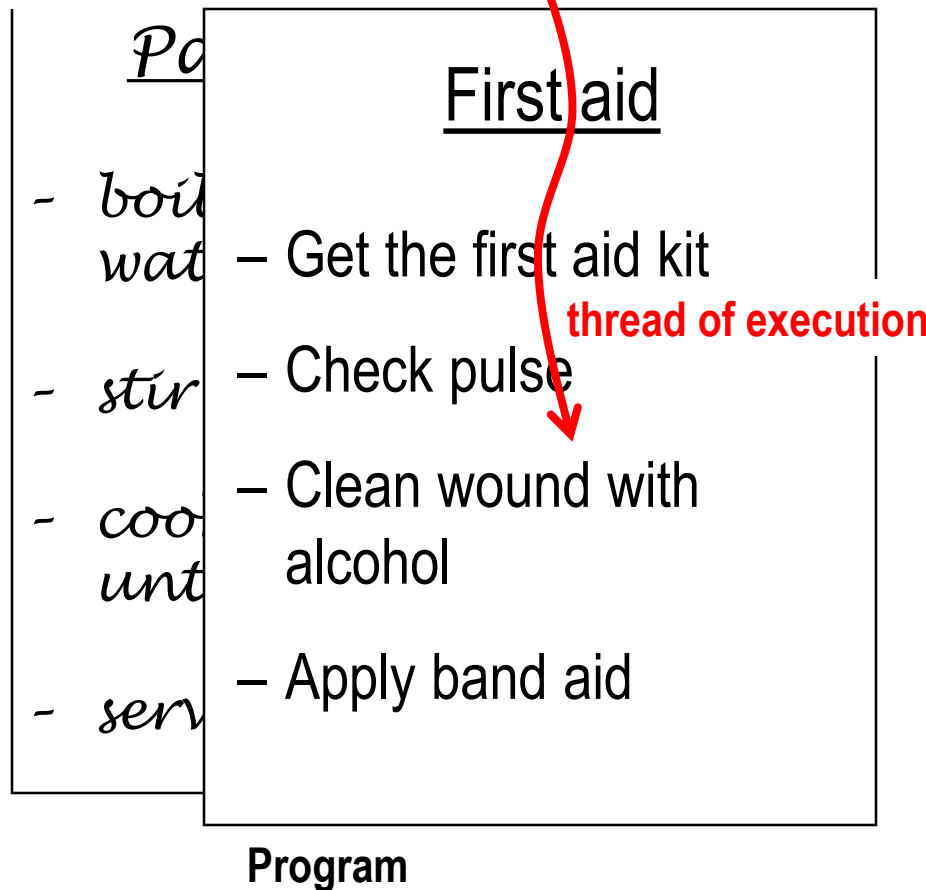
input data

Process

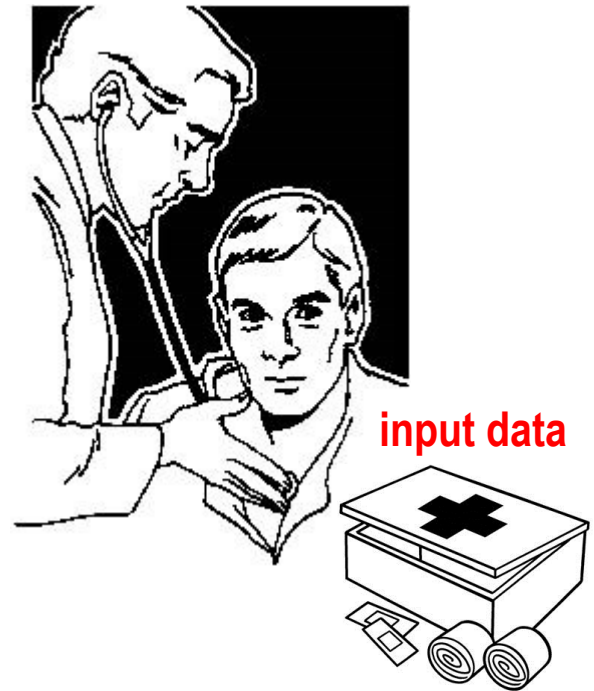
# a. Processes (Summary)

What is a process?

- It can be **interrupted** to let the CPU execute a higher-priority process



**CPU (changes hat to “doctor”)**

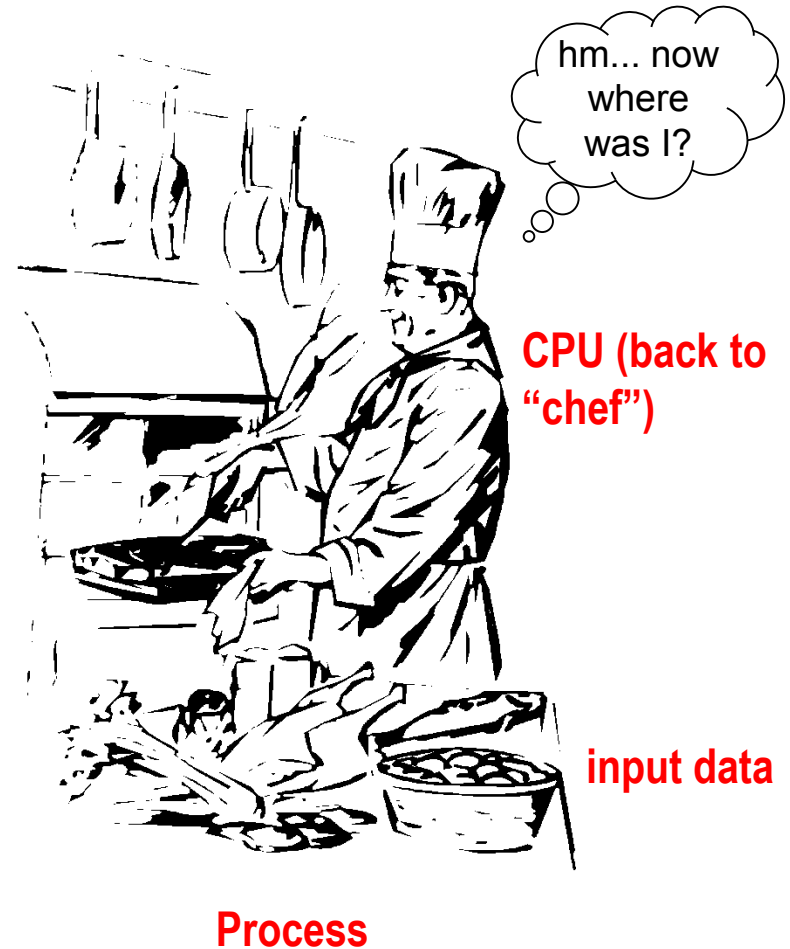
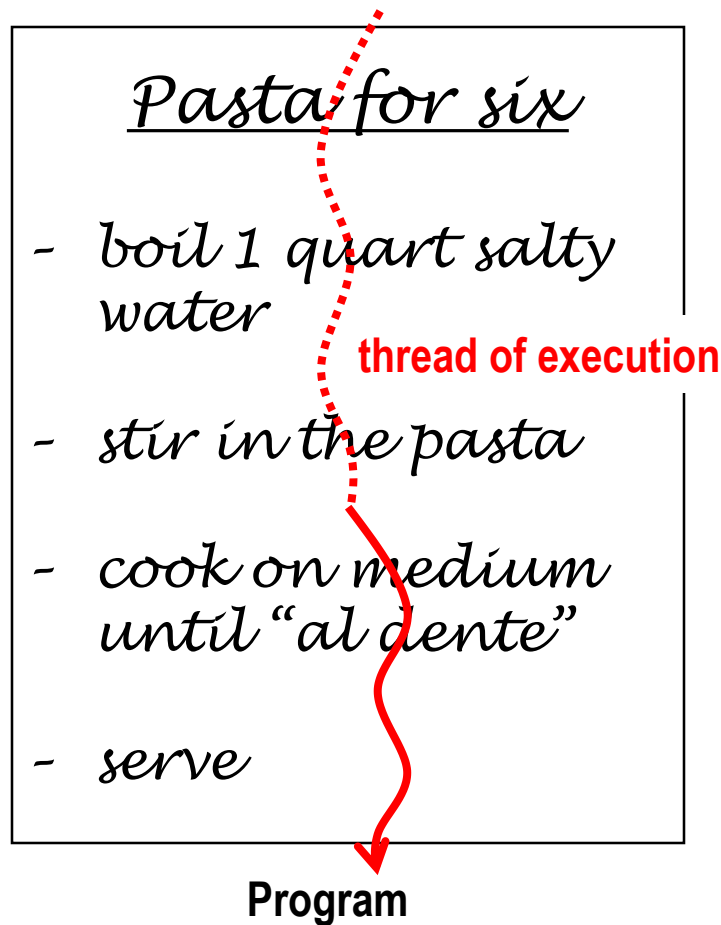


**Process**

# a. Processes (Summary)

What is a process?

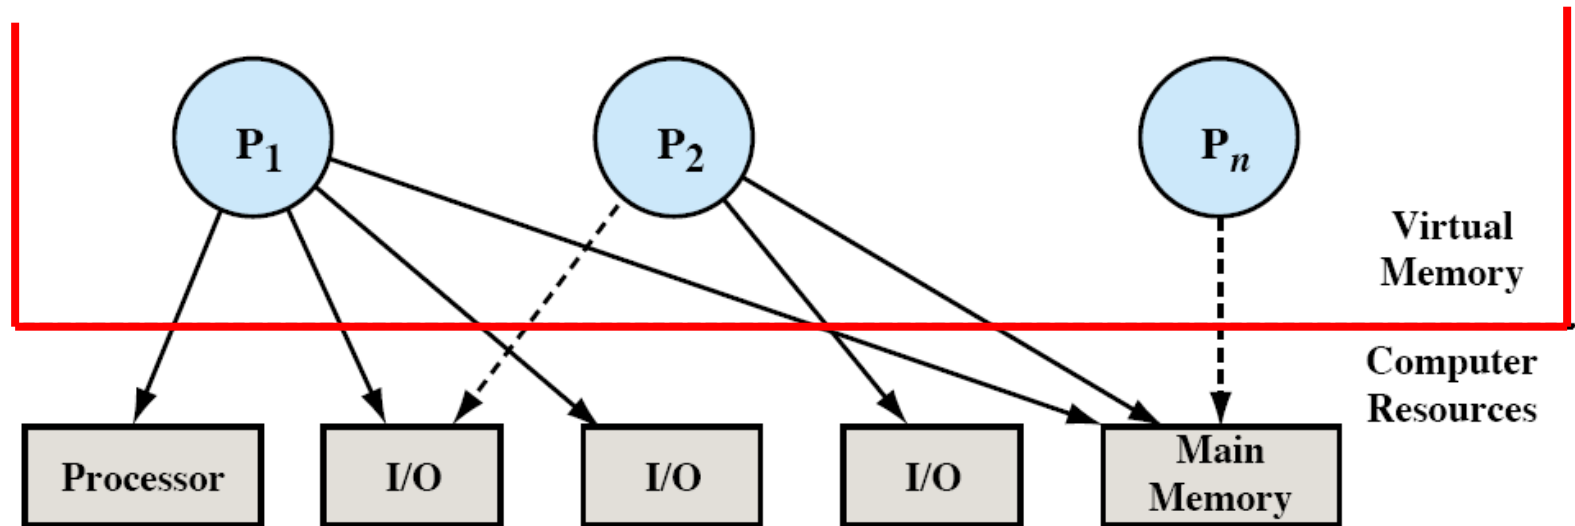
- . . . and then **resumed** exactly where the **CPU left off**



# a. Processes (Summary)

## What is a process?

- The OS has to **multiplex** resources to the processes
  - ✓ a number of processes have been created
  - ✓ each process during the course of its execution needs access to system resources: **CPU, main memory, I/O devices**



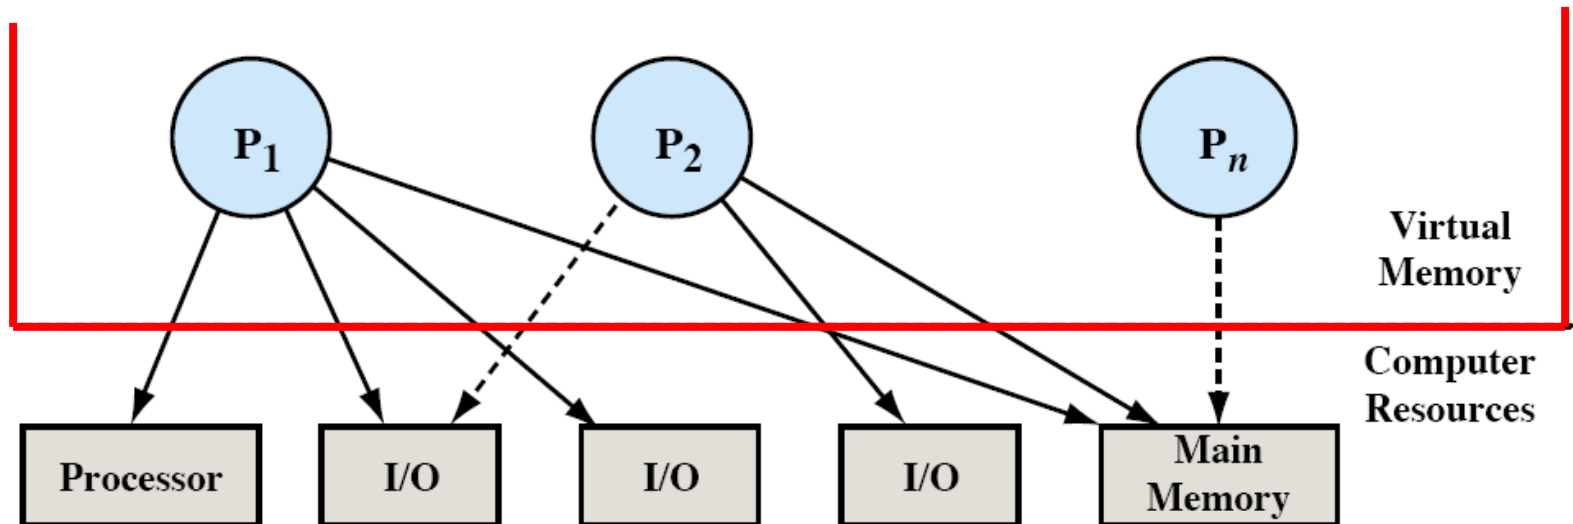
Stallings, W. *Operating Systems: Internals and Design Principles*.

Resource allocation for processes (one snapshot in time)

# a. Processes (Summary)

## What is a process?

- $P_1$  is **running**; at least part of the process is in **main memory**, and it has control of **two** I/O devices
- $P_2$  is also in main memory but is **blocked** waiting for an I/O device allocated to  $P_1$
- $P_n$  has been **swapped out** and is therefore **suspended**



Stallings, W. *Operating Systems: Internals and Design Principles*.

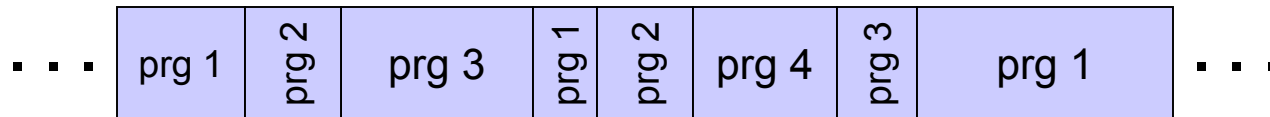
Resource allocation for processes (one snapshot in time)



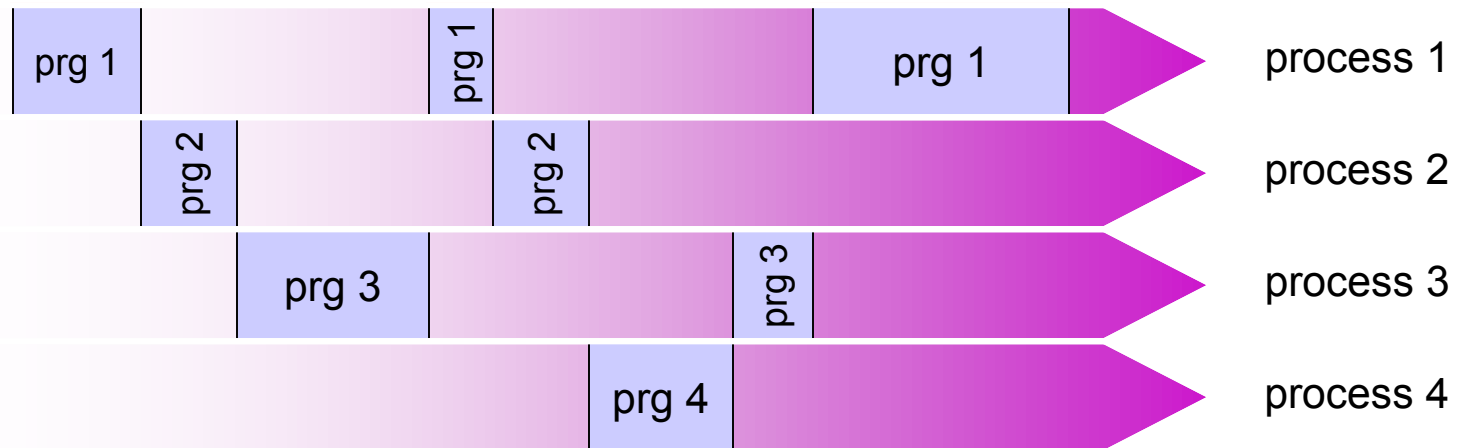
# a. Processes (Summary)

What is a process?

- Multitasking can be conveniently described in terms of multiple processes running in (pseudo) parallel



(a) Multitasking from the CPU's viewpoint



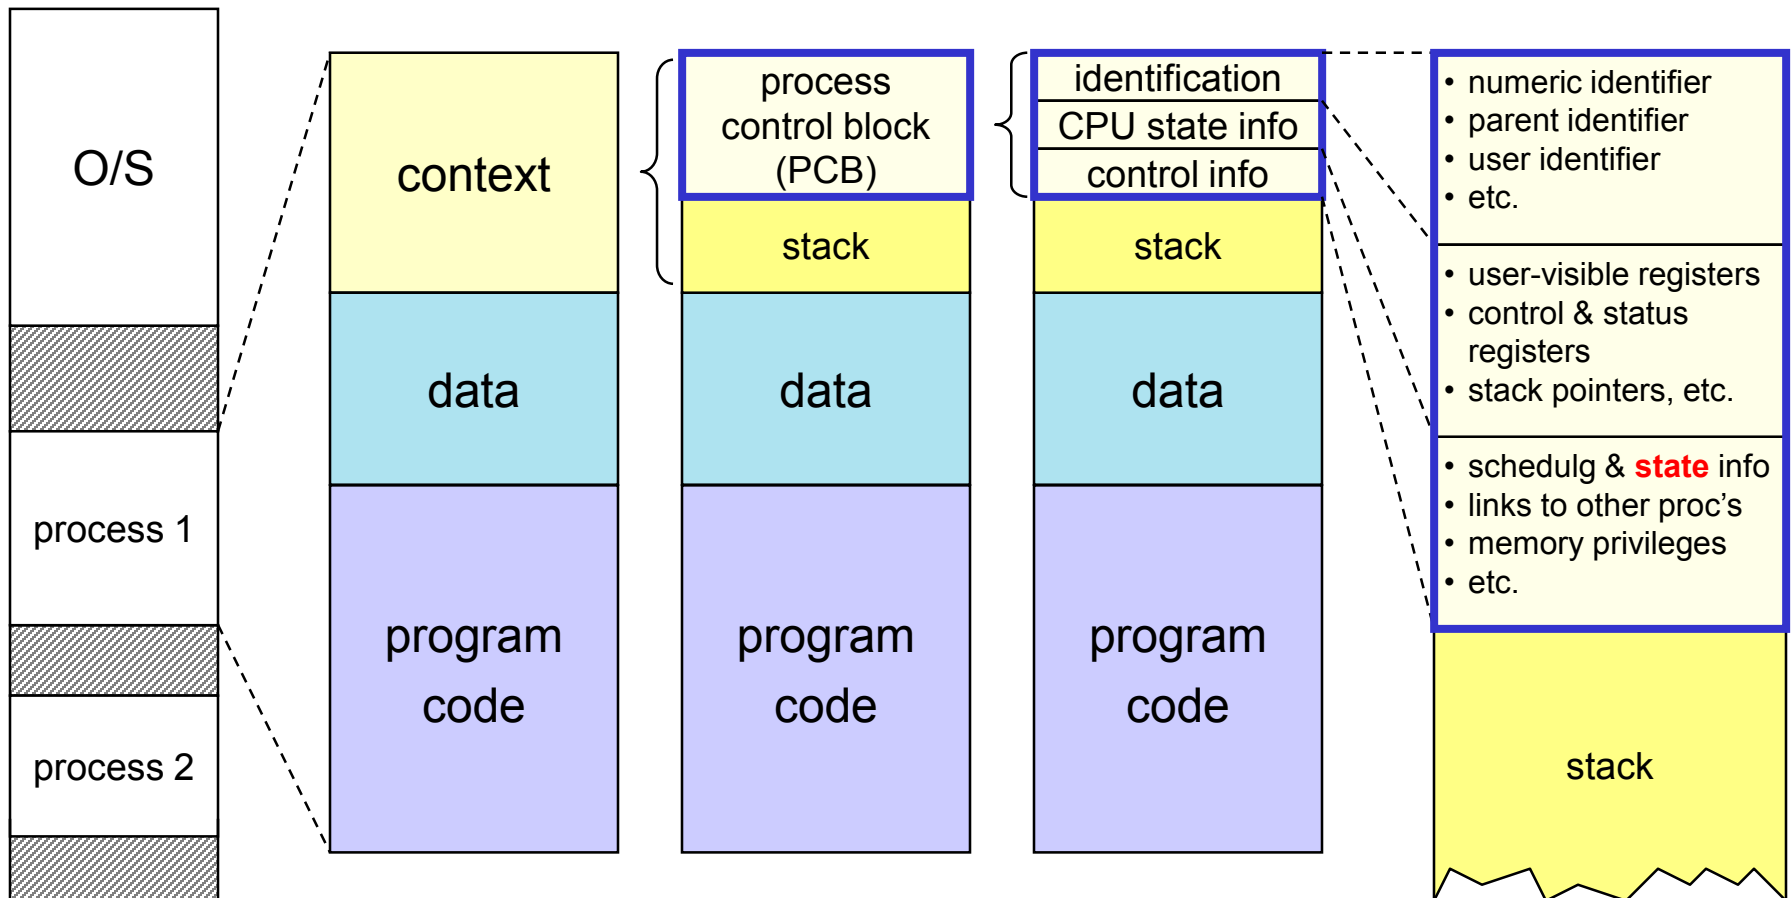
(b) Multitasking from the processes' viewpoint = 4 virtual program counters

**Pseudo-parallelism in multitasking**

# a. Processes (Summary)

## Process description

- In the process table, the OS keeps one structure per process, the *Process Control Block (PCB)*, containing:



Illustrative contents of a process image in (virtual) memory

# a. Processes (Summary)

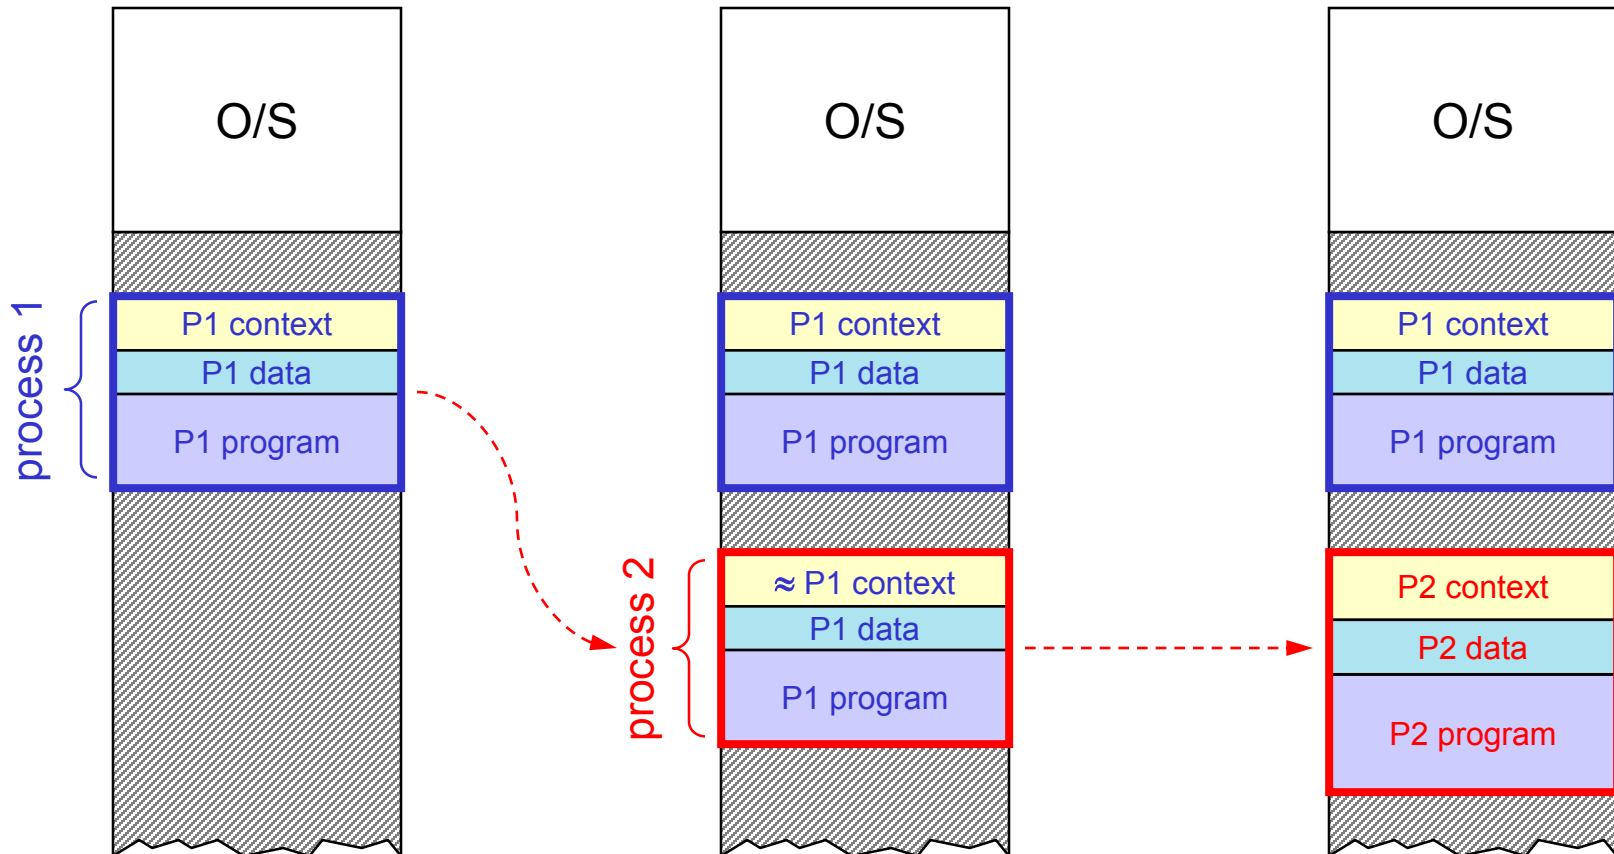
## Process control

1. Clone **child** process

✓ `pid = fork()`

2. Replace child's image

✓ `execve(name, ...)`



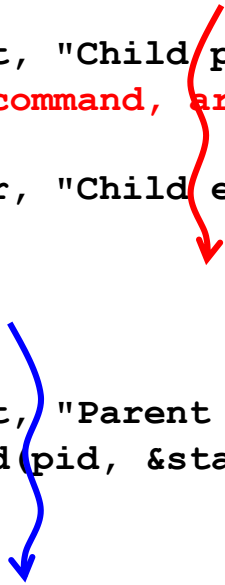
# a. Processes (Summary)

## Process control

```
...
int main(...)                // the same program will be executed twice, but differently (not concurrently)
{
    ...
    if ((pid = fork()) == 0) // create a process:
    {                          // we are in the child process
        fprintf(stdout, "Child pid: %i\n", getpid());
        err = execvp(command, arguments); // execute child process

        fprintf(stderr, "Child error: %i\n", errno);
        exit(err);
    }
    else if (pid > 0)        // we are in the parent process
    {                          // and pid is the child's ID
        fprintf(stdout, "Parent pid: %i\n", getpid());
        pid2 = waitpid(pid, &status, 0); // wait for child process
        ...
    }
    ...

    return 0;
}
```



Implementing a shell command interpreter  
by **process spawning**

# a. Processes (Summary)

## Process control

- What events trigger the OS to switch processes?
  - ✓ **interrupts** — external, asynchronous events, independent of the currently executed process instructions
    - **Clock interrupt** → OS checks time and may block process
    - **I/O interrupt** → data has come, OS may unblock process
    - **Memory fault** → OS may block process that must wait for a **missing page** in memory to be swapped in

# a. Processes (Summary)

## Process control

### Traps

✓ **exceptions** — internal, involuntary synchronous events  
caused by instructions

→ OS may **terminate** or **recover** process

✓ **system calls** — internal, voluntary synchronous events  
calling a specific OS service

→ after service completed, OS may either **resume** or  
**block** the calling process, depending on I/O, priorities, etc.

- a. Processes (Summary)
- b. Threads (Summary)
  - ✓ Separation of resource ownership and execution
  - ✓ Same old throughput story
  - ✓ Implementation of threads
  - ✓ Practical uses of multithreading
- c. Concurrency
- d. Deadlock

# Threads

## Separation of resource ownership and execution

➤ In fact, a process embodies two independent concepts

1. resource ownership
2. execution & scheduling

### 1. Resource ownership

- ✓ a process is allocated address space to hold the image, and is granted control of I/O devices and files
- ✓ the OS **prevents interference** among processes while they make use of resources (**multiplexing**)

### 2. Execution & scheduling

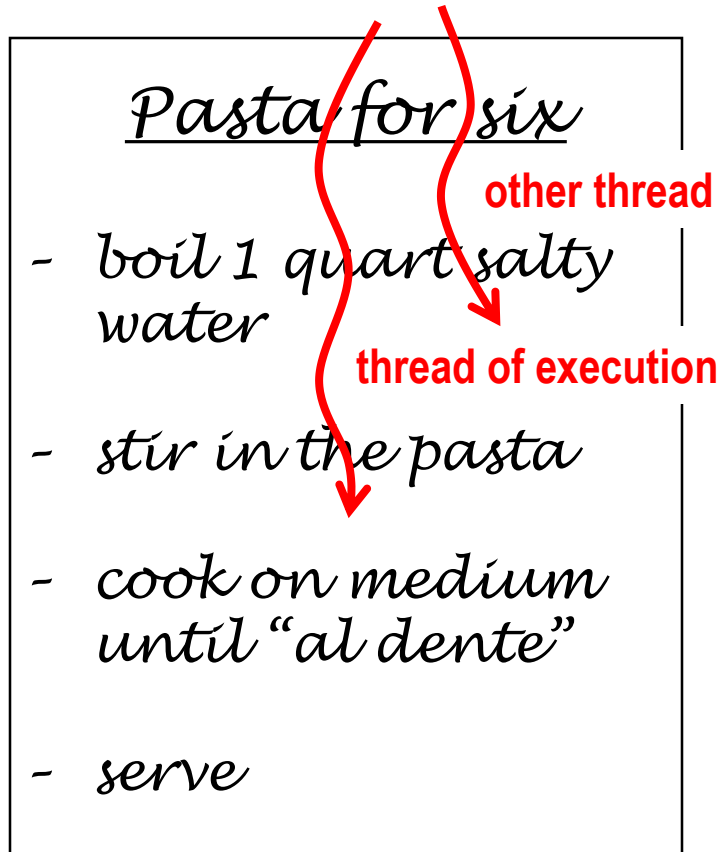
- ✓ a process follows an **execution path** through a program
- ✓ it has an execution state and is scheduled for dispatching



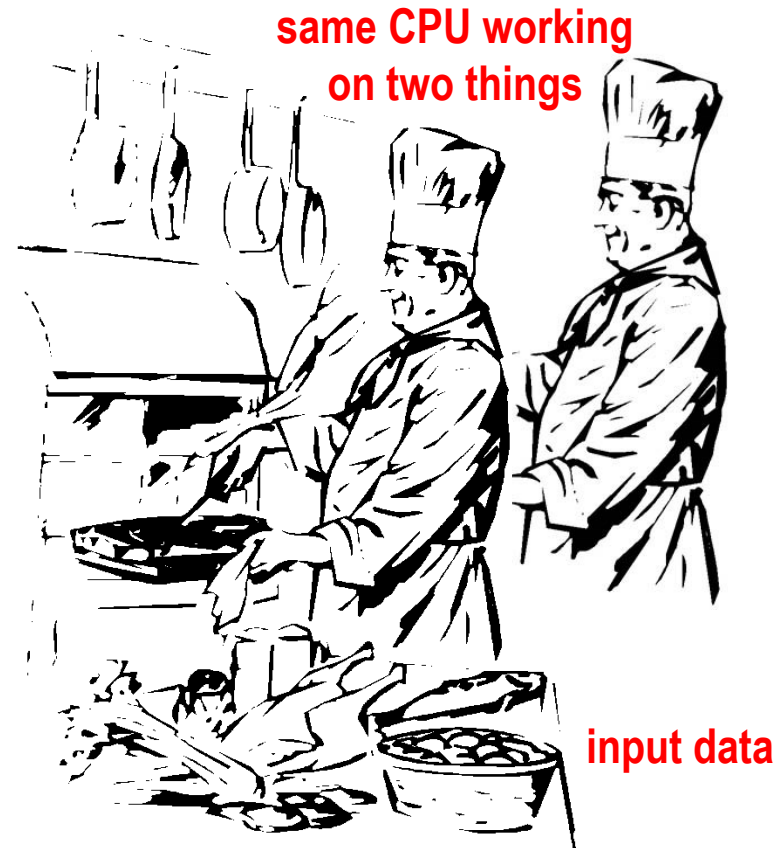
# Threads

**Separation** of resource ownership and execution

- The execution part is a “**thread**” that can be multiplied



Program



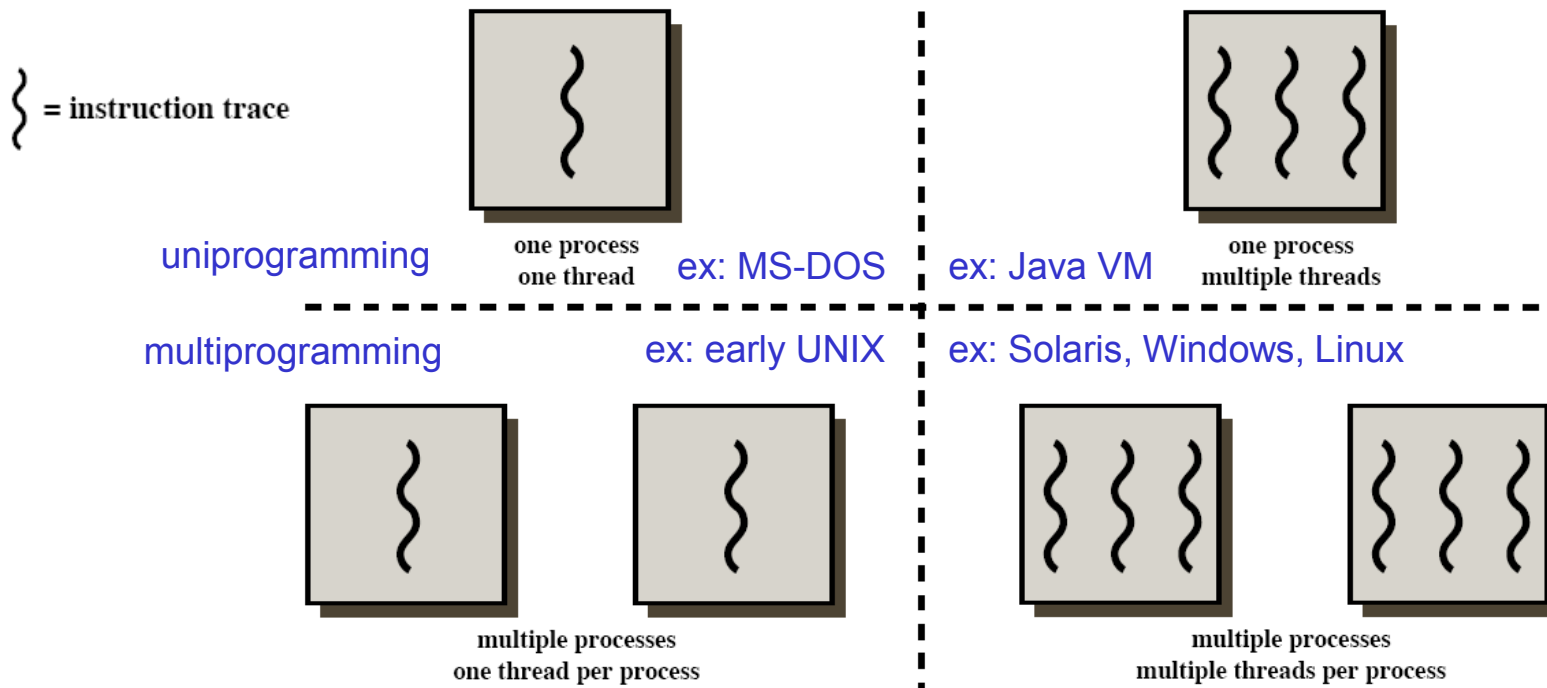
Process

# Threads

## Separation of resource ownership and execution

### ➤ Multithreading

- ✓ refers to the ability of an OS to support **multiple threads** of execution within a single process



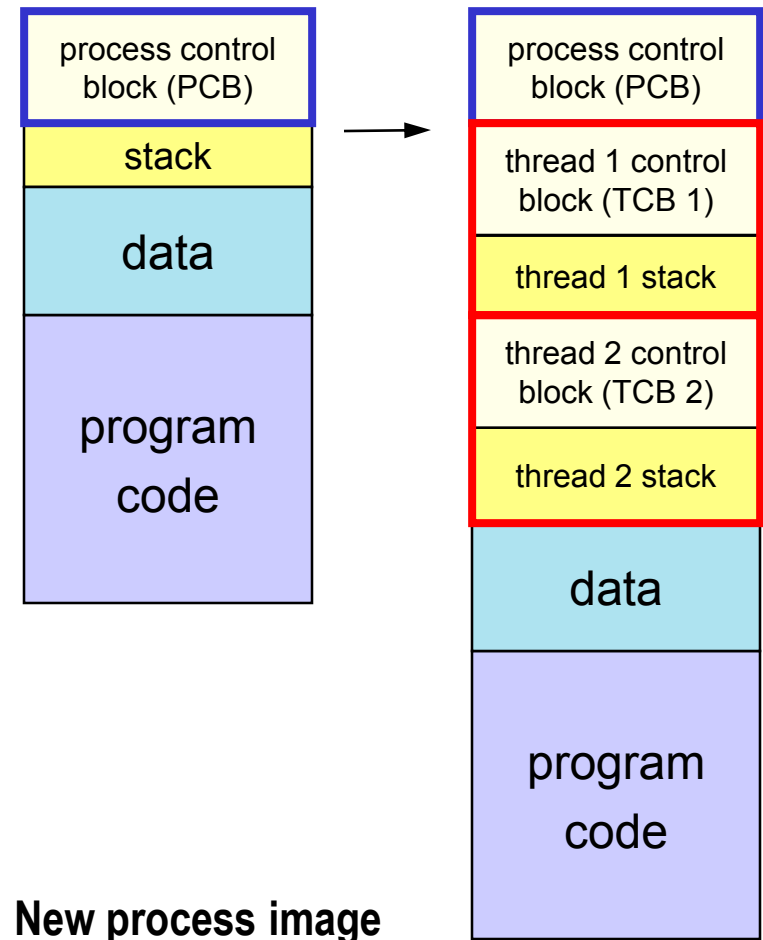
Process-thread relationships

## b. Threads (Summary)

### Separation of resource ownership and execution

➤ Multithreading requires **changes** in the process description model

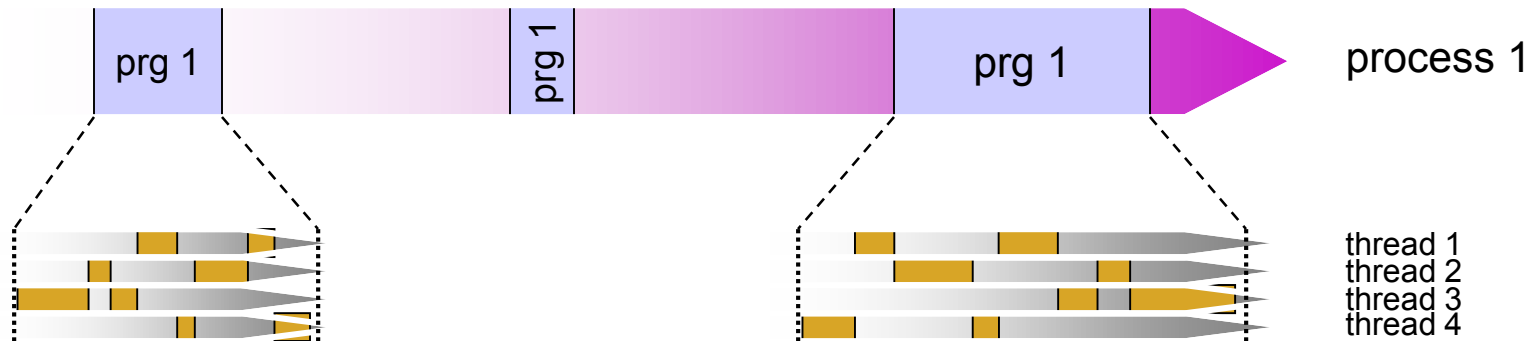
- ✓ each thread of execution receives its **own** control block and stack
  - own execution state (“Running”, “Blocked”, etc.)
  - own copy of CPU registers
  - own execution history (stack)
- ✓ the process keeps a **global control block** listing resources currently used



## b. Threads (Summary)

Same old throughput story

- The same multitasking idea applies in multithreading
  - ✓ **multithreading** is basically the same as **multitasking** at a finer level of temporal resolution (and within the same address space)
  - ✓ the same illusion of parallelism is achieved at a **finer grain**

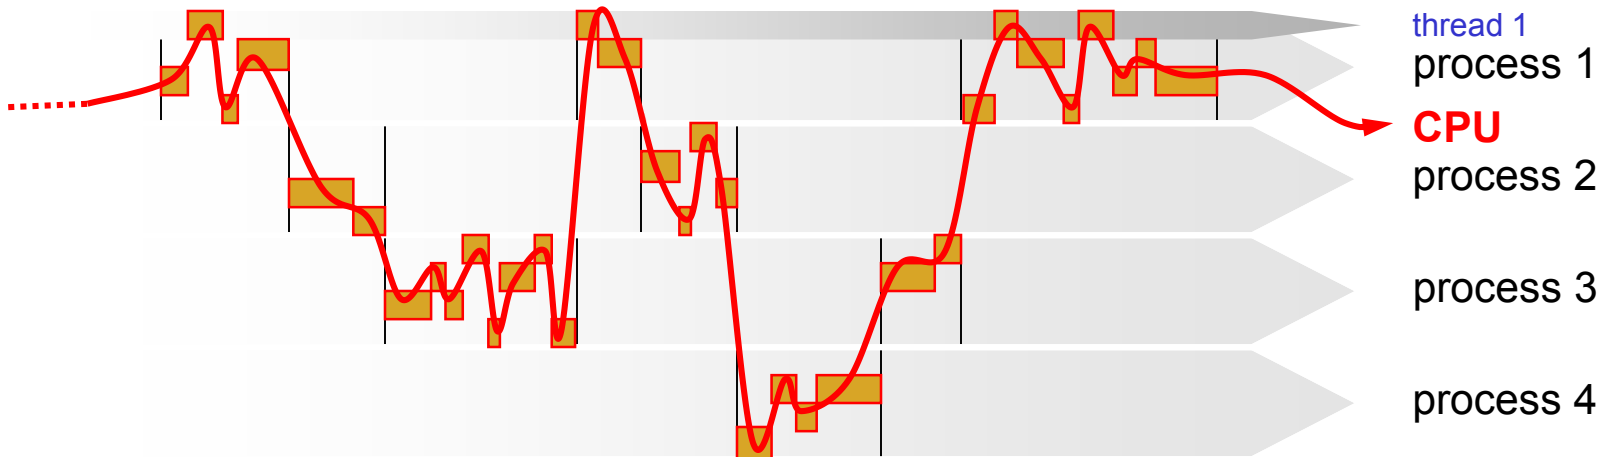


**Multithreading**

## b. Threads (Summary)

Same old throughput story

- **The same multitasking idea applies in multithreading**
  - ✓ in a single-processor system, there is still only one CPU going through all the threads of all the processes



**Multithreading**

## b. Threads (Summary)

### Same old throughput story

- Benefits of multithreading compared to multitasking
  - ✓ it takes **less time** to **create** a new thread than a new process
  - ✓ it takes **less time** to **terminate** a thread than a process
  - ✓ it takes **less time** to **switch** between two threads within the same process than between two processes
  - ✓ threads within the same process **share memory and files**, therefore they can communicate with each other without having to **invoke the kernel**
  - ✓ for these reasons, threads are sometimes called “**lightweight processes**”
- *if an application should be implemented as a set of related executions, **it is far more efficient to use threads than processes***

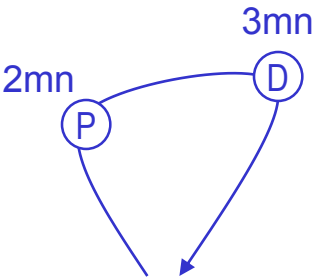
## b. Threads (Summary)

### Practical uses of multithreading

#### ➤ Illustration: two shopping scenarios

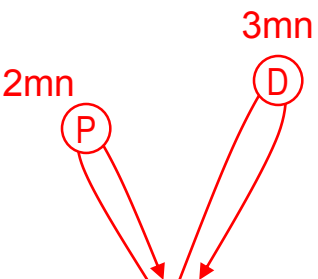
##### ✓ Single-threaded shopping

- you are in the grocery store
- first you go to produce and grab salad and apples, then you go to dairy and grab milk, butter and cheese
- it took you about  $1mn \times 5 \text{ items} = 5mn$



##### ✓ Multithreaded shopping

- you take your **two kids** with you to the grocery store
- you send them off in **two directions** with **two missions**, one toward produce, one toward dairy
- you wait for their return for a **maximum duration** of about  $1mn \times 3 \text{ items} = 3mn$



## b. Threads (Summary)

### Practical uses of multithreading

#### ➤ Examples of real-world multithreaded applications

##### ✓ Web client (browser)

- must download page components (images, styles, etc.) simultaneously; cannot wait for each image in series

##### ✓ Web server

- must serve pages to hundreds of Web clients simultaneously; cannot process requests one by one

##### ✓ Word processor, spreadsheet

- provides uninterrupted GUI service to the user while reformatting or saving the document in the background

→ *again, same principles as time-sharing (illusion of **interactivity** while performing other tasks), this time inside the same process*



- a. Processes (Summary)
- b. Threads (Summary)
- c. Concurrency
  - ✓ Types of process interaction
  - ✓ Race conditions & critical sections
  - ✓ Mutual exclusion by busy waiting
  - ✓ Mutual exclusion & synchronization
- d. Deadlock

## c. Concurrency

### Types of process interaction

- Concurrency refers to any form of interaction among processes or threads
  - ✓ concurrency is a fundamental part of OS design
  - ✓ concurrency includes
    - **communication** among processes/threads
    - **sharing** of, and **competition** for, system resources
    - **cooperative** processing of shared data
    - **synchronization** of process/thread activities
    - organized CPU scheduling
    - solving **deadlock** and **starvation** problems

## c. Concurrency

### Types of process interaction

- Concurrency arises in the same way at different levels of execution streams

Software  
view

- ✓ **multiprogramming** — interaction between multiple processes running on one CPU (pseudo-parallelism)
- ✓ **multithreading** — interaction between multiple threads running in one process

Hardware  
view

- ✓ **multiprocessors** — interaction between multiple CPUs running multiple processes/threads (real parallelism)
- ✓ **multicomputers** — interaction between multiple computers running **distributed** processes/threads

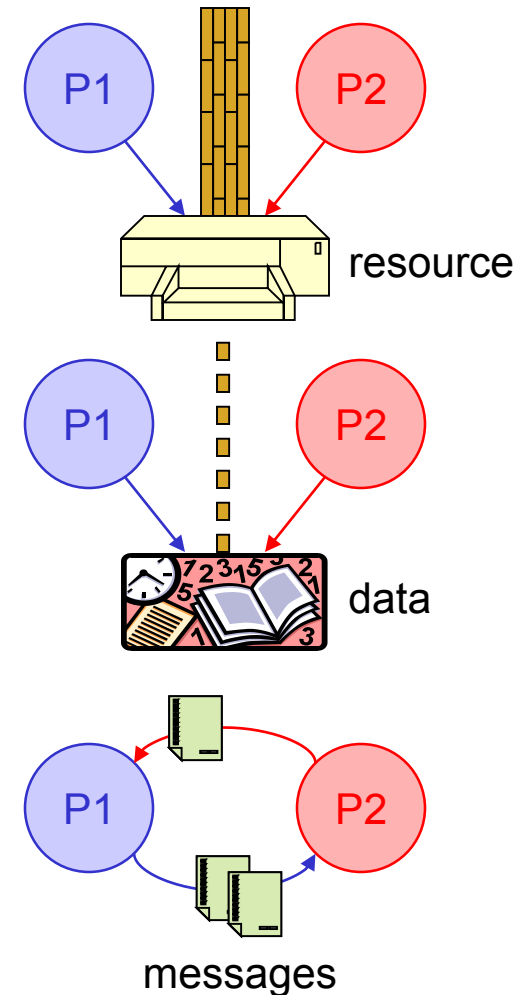
→ *the principles of concurrency are basically the same in all of these categories*

## c. Concurrency

### Types of process interaction

#### ➤ Whether processes or threads: three basic interactions

- ✓ **processes unaware of each other** — they must use **shared resources** independently, without interfering, and leave them intact for the others
- ✓ **processes indirectly aware of each other** — they work on **common data** and build some result together via the data
- ✓ **processes directly aware of each other** — they cooperate by **communicating**, e.g., exchanging messages



## c. Concurrency

### Race conditions & critical sections

- **Race condition:** a situation in which multiple threads or processes read and write a shared data item and the final result depends on the **relative timing** of their execution.
  - ✓ there is a “**race condition**” if the outcome depends on the order of the execution
- **Critical section (region):** a section of code within a process that requires access to **shared resources** and that must not be executed while another process is in a corresponding section of code.
- **Mutual exclusion:** the requirement that when one process is in a critical section that accesses shared resources, **no other process** may be in a critical section that accesses **any of those shared resources**.

## c. Concurrency

### Race conditions & critical sections

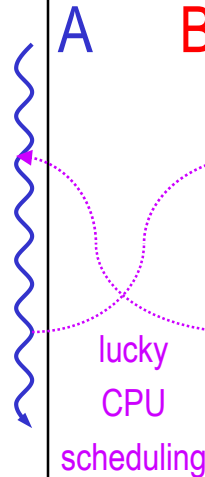
#### ➤ Significant race conditions in I/O & variable sharing

```
char chin, chout;

void echo()
{
    do {
        1 chin = getchar();
        2 chout = chin;
        3 putchar(chout);
    }
    while (...);
}
```

```
> ./echo
Hello world!
Hello world!
```

Single-threaded echo



```
char chin, chout;

void echo()
{
    do {
        4 chin = getchar();
        5 chout = chin;
        6 putchar(chout);
    }
    while (...);
}
```

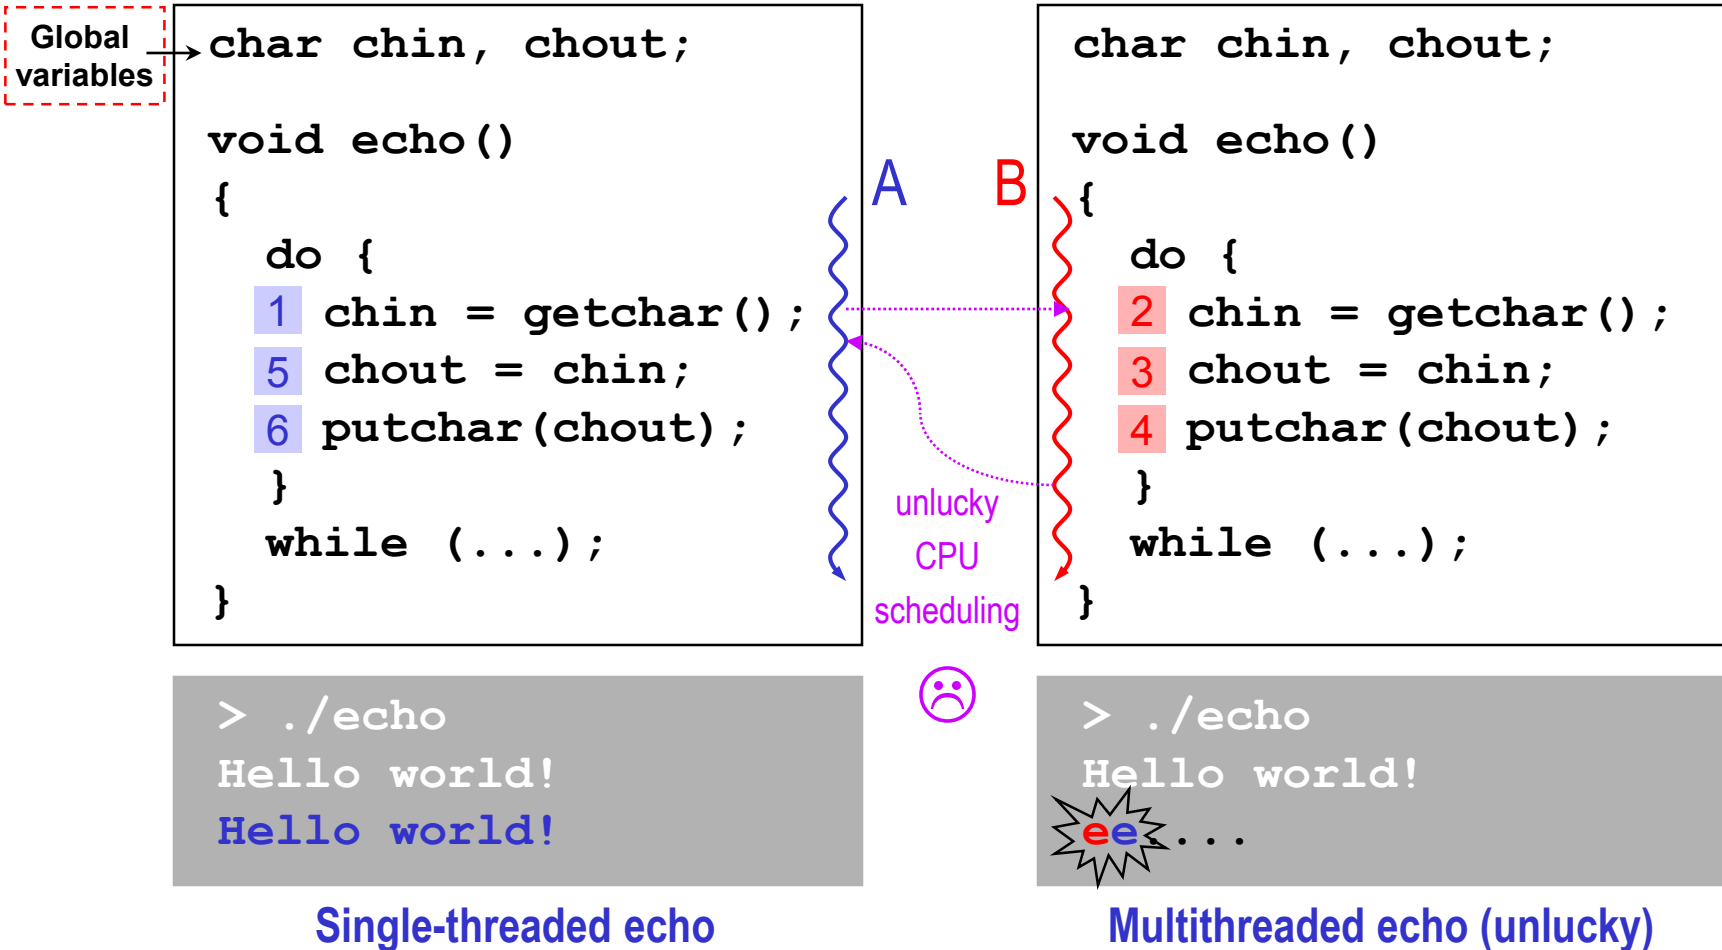
```
> ./echo
Hello world!
Hello world!
```

Multithreaded echo (lucky)

# c. Concurrency

## Race conditions & critical sections

### ➤ Significant race conditions in I/O & variable sharing



# c. Concurrency

## Race conditions & critical sections

### ➤ Significant race conditions in I/O & variable sharing

changed to local variables →

```
void echo()
{
    char chin, chout;

    do {
        1 chin = getchar();
        5 chout = chin;
        6 putchar(chout);
    }
    while (...);
}
```

```
> ./echo
Hello world!
Hello world!
```

Single-threaded echo

A B

```
void echo()
{
    char chin, chout;

    do {
        2 chin = getchar();
        3 chout = chin;
        4 putchar(chout);
    }
    while (...);
}
```

unlucky  
CPU  
scheduling



```
> ./echo
Hello world!
eH...
```

Multithreaded echo (unlucky)



## c. Concurrency

### Race conditions & critical sections

#### ➤ Significant race conditions in I/O & variable sharing

- ✓ note that, in this case, replacing the **global** variables with **local** variables did not solve the problem
- ✓ we actually had two race conditions here:
  - one race condition in the shared variables and the order of value assignment
  - another race condition in the shared output stream: which thread is going to write to output first (this race persisted even after making the variables local to each thread)
- generally, problematic race conditions may occur whenever **resources and/or data are shared** (by processes **unaware** of each other or processes **indirectly aware** of each other)

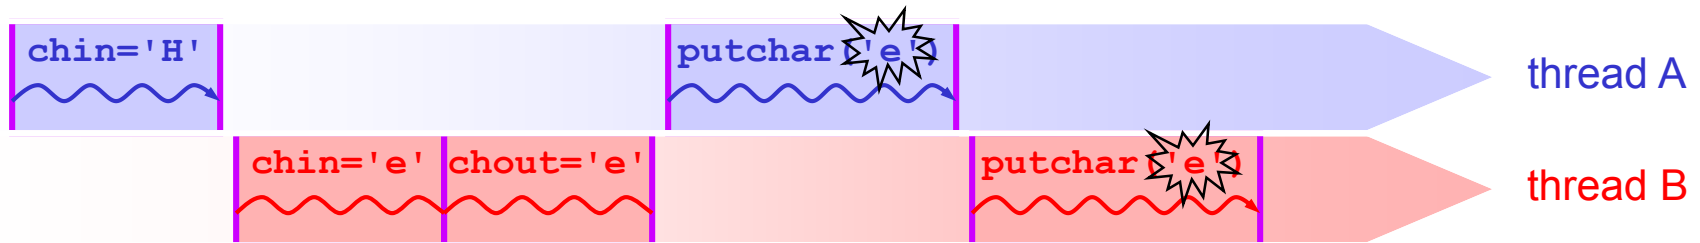
**How to avoid race conditions?**

## c. Concurrency

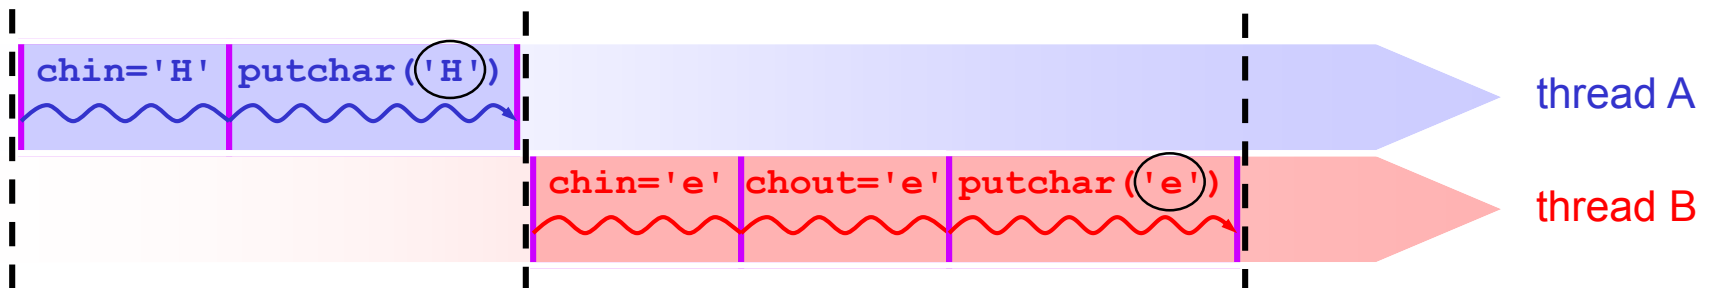
### Race conditions & critical sections

#### ➤ How to avoid race conditions?

- ✓ find a way to keep the instructions together
- ✓ this means actually. . . reverting from too much interleaving and going back to “**indivisible**” blocks of execution!!



(a) too much interleaving may create race conditions



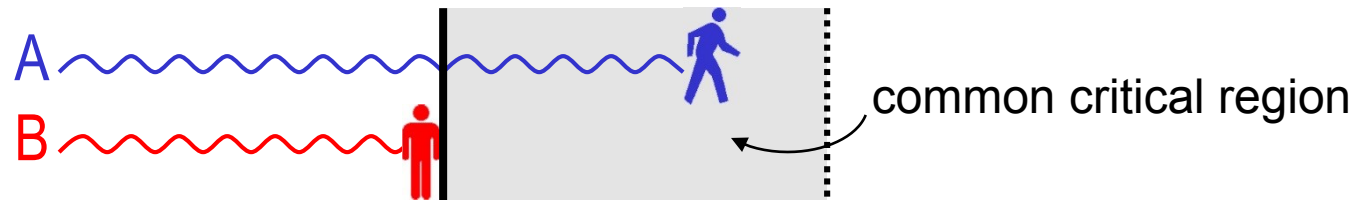
(b) keeping “indivisible” blocks of execution avoids race conditions

## c. Concurrency

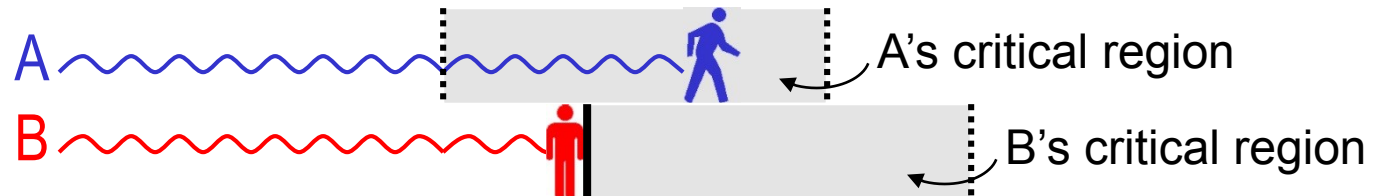
### Race conditions & critical sections

- The “indivisible” execution blocks are critical sections

✓ *a critical region is a section of code that may be executed by only one process or thread at a time*



- ✓ although it does not have to be the same region of memory or section of program in both processes



→ *but physically different or not, what matters is that these regions cannot be **interleaved** or **executed in parallel** (pseudo or real)*


## c. Concurrency

### Race conditions & critical regions

#### ➤ We need mutual exclusion from critical regions

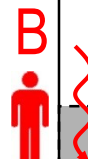
- ✓ critical regions can be protected from concurrent access by padding them with **entrance and exit gates**:
  - ✓ a thread must try to **check in**, then it must **check out**

```
void echo()  
{  
    char chin, chout;  
    do {  
        enter critical region?  
        chin = getchar();  
        chout = chin;  
        putchar(chout);  
        exit critical region  
    }  
    while (...);  
}
```



A blue wavy line labeled 'A' represents thread A's attempt to enter the critical region. A blue stick figure is shown at the bottom of the wavy line, indicating the thread is currently inside the critical region.

```
void echo()  
{  
    char chin, chout;  
    do {  
        enter critical region?  
        chin = getchar();  
        chout = chin;  
        putchar(chout);  
        exit critical region  
    }  
    while (...);  
}
```



A red wavy line labeled 'B' represents thread B's attempt to enter the critical region. A red stick figure is shown at the top of the wavy line, indicating the thread is blocked from entering because the critical region is already occupied by thread A.

## c. Concurrency

### Race conditions & critical regions

#### Chart of mutual exclusion

1. **mutual exclusion inside** — only **one process** at a time may be allowed in a critical region
2. **no exclusion outside** — a process stalled in a **non-critical region** may not exclude other processes from their critical regions
3. **no indefinite occupation** — a critical region may be only occupied for a finite amount of time

## c. Concurrency

### Race conditions & critical regions

#### Chart of mutual exclusion (cont'd)

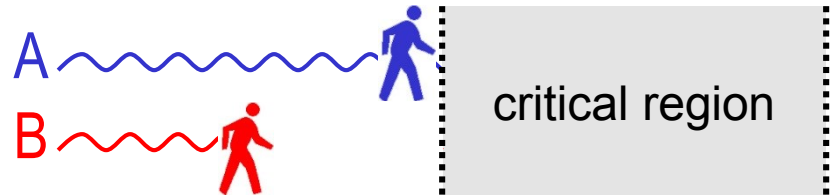
4. **no indefinite delay** — a process may be only excluded for a finite amount of time (no **deadlock** or **starvation**)
5. **no delay when about to enter** — a critical region free of access may be entered **immediately** by a process
6. **nondeterministic scheduling** — no assumption should be made about the relative speeds of processes

## c. Concurrency

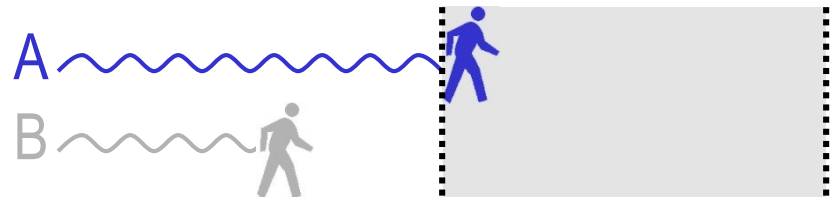
Mutual exclusion by **busy waiting**

### ➤ Disabling **hardware** interrupts

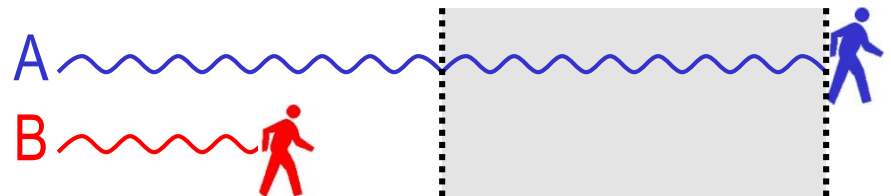
1. thread A reaches the gate to the critical region (CR) before B



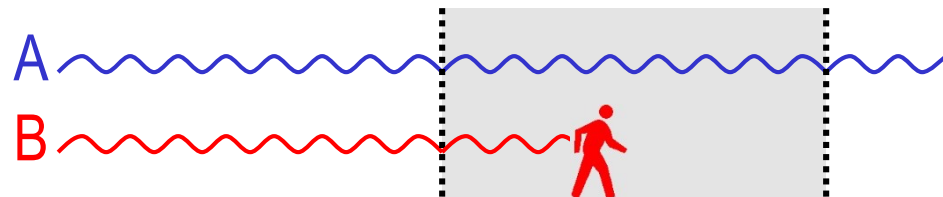
2. as soon as A enters CR, it **disables all interrupts**, thus B cannot be scheduled



3. as soon as A exits CR, it **reenables interrupts**; B can be scheduled again



4. thread B enters CR





## c. Concurrency

### Mutual exclusion by **busy waiting**

#### ➤ Disabling **hardware** interrupts 🖱

- ✓ it works, but is foolish
- ✓ what guarantees that the user process is going to ever exit the critical region?
- ✓ meanwhile, the **CPU cannot interleave any other task**, even unrelated to this race condition
- ✓ the critical region becomes one **physically** indivisible block, not logically.
- ✓ also, this is not working in multi-processors

```
void echo()  
{  
    char chin, chout;  
    do {  
        disable hardware interrupts  
        chin = getchar();  
        chout = chin;  
        putchar(chout);  
        reenable hardware interrupts  
    }  
    while (...);  
}
```

## c. Concurrency

### Mutual exclusion by **busy waiting**

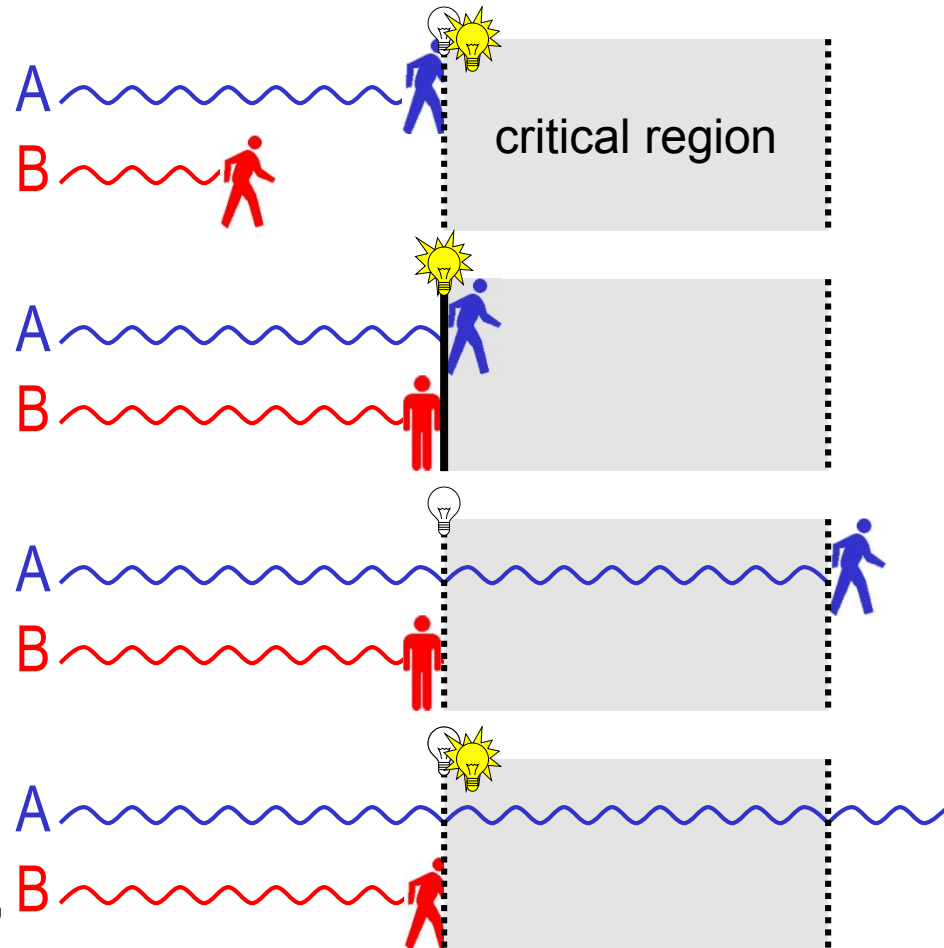
#### ➤ **Indivisible lock variable** 👍

1. thread A reaches CR and finds the lock at 0 and sets it in **one shot**, then enters

1.1' even if B comes right behind A, it will find that the lock is already at 1

2. thread A exits CR, then resets lock to 0

3. thread B finds the lock at 0 and sets it to 1 in **one shot**, just before entering CR



## c. Concurrency

### Mutual exclusion by busy waiting

#### ➤ Indivisible lock variable 👍

- ✓ the indivisibility of the “test-lock-and-set-lock” operation can be implemented with a single hardware instruction **TSL**

```
enter_region:
TSL REGISTER, LOCK | copy lock to register and set lock to 1 (*)
CMP REGISTER, #0    | was lock zero?
JNE enter_region    | if it was non zero, lock was set, so loop
RET                 | return to caller; critical region entered
```

```
leave_region:
MOVE LOCK, #0       | store a 0 in lock
RET                 | return to caller
```

```
void echo()
{
    char chin, chout;
    do {
        test-and-set-lock
        chin = getchar();
        chout = chin;
        putchar(chout);
        set lock off
    }
    while (...);
}
```

## c. Concurrency

Mutual exclusion by **busy waiting**

### ➤ Summary of mutual exclusion implementations

#### ➤ Disabling hardware interrupts

- NO: race condition avoided, but can crash the system!

#### ➤ Indivisible lock variable (TSL)

- YES: works, but requires special hardware instruction
  - It is the basis for mutexes

#### ➤ Other implementations:

##### ➤ Simple lock variable (unprotected)

- NO: still suffers from race condition

##### ➤ Peterson's no-TSL, no-alternation

- YES: works in pure software, but processing overhead

## c. Concurrency

### Mutual exclusion by **busy waiting**

#### ➤ **Problem: all implementations rely on busy waiting**

- ✓ “busy waiting” means that any process/thread (A, B, ...) must continuously **execute a tight loop** until some condition changes
- ✓ busy waiting is bad:
  - **waste of CPU time** — the busy process is not doing anything useful, yet remains “Ready” instead of “Blocked”
  - **paradox of inversed priority** — by looping indefinitely, a higher-priority process B may starve a lower-priority process A, thus **preventing A from exiting CR** and . . . keeping B in the wait! (the CPU is too busy executing an empty loop, i.e. B is working against its own interest)

→ *we need for the waiting process to truly block, not keep idling*

## ✓ Mutual exclusion & synchronization

- Mutexes
- Semaphores
- Monitors

## c. Concurrency

### Mutual exclusion & synchronization — **mutexes**

#### ➤ Indivisible blocking lock = **mutex**

- ✓ a mutex is a safe lock variable with **blocking**, instead of tight looping
- ✓ if **TSL** returns 1, then voluntarily yield the CPU to another thread

```
mutex_lock:
  TSL REGISTER,MUTEX | copy mutex to register and set mutex to 1
  CMP REGISTER,#0    | was mutex zero?
  JZE ok              | if it was zero, mutex was unlocked, so return
  CALL thread_yield   | mutex is busy; schedule another thread
  JMP mutex_lock      | try again later
ok: RET               | return to caller; critical region entered
```

```
mutex_unlock:
  MOVE MUTEX,#0       | store a 0 in mutex
  RET                  | return to caller
```

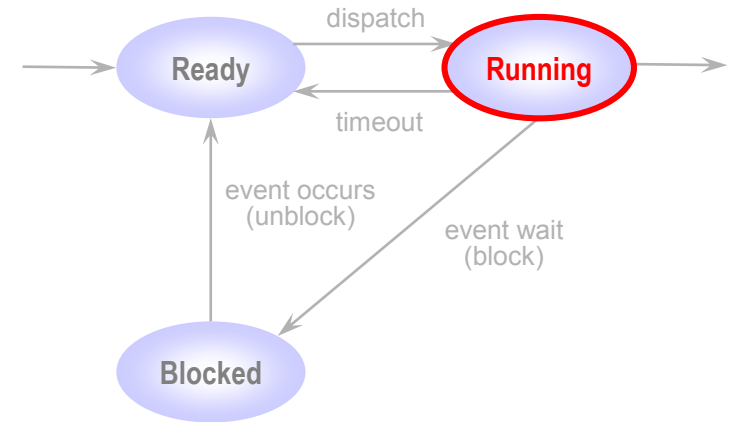
```
void echo()
{
  char chin, chout;
  do {
    test-and-set-lock or BLOCK
    chin = getchar();
    chout = chin;
    putchar(chout);
    set lock off
  }
  while (...);
}
```

## c. Concurrency

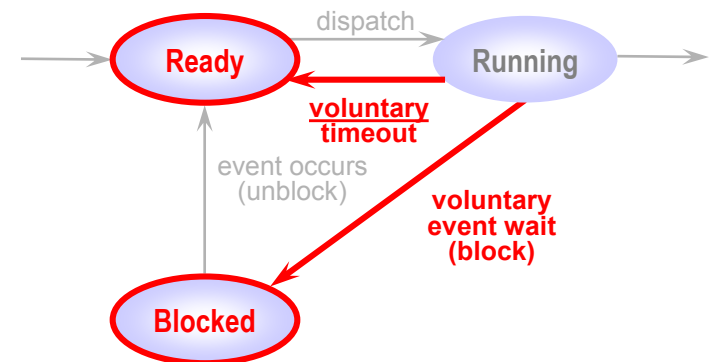
### Mutual exclusion & synchronization — **mutexes**

#### ➤ Difference between **busy waiting** and **blocked**

- ✓ in busy waiting, the PC (Program Counter) is always **looping** (increment & jump back)
- ✓ it can be preemptively interrupted but will loop again tightly whenever rescheduled → *tight polling*



- ✓ when blocked, the process's PC stalls after executing a “**yield**” call
- ✓ either the process is only timed out, thus it is “**Ready**” to loop-and-yield again → *sparse polling*
- ✓ or it is truly “**Blocked**” and put in event queue → *condition waiting*





## c. Concurrency

### Mutual exclusion & synchronization — Semaphores

- Special integer variable with only **three possible operations**:
  - ✓ **Initialisation** (to a non-negative number (0 or 1 for binary sem))
  - ✓ **wait**, which decrements the value
    - ✓ If the value is then **negative**, the process calling wait becomes **blocked**
  - ✓ **signal**, which increments the value
    - ✓ If the value is then  $\leq 0$ , a process which has previously been blocked by this semaphore becomes unblocked
  - ✓ Important: wait and signal should be implemented as **atomic (indivisible) primitives**

## c. Concurrency

### Mutual exclusion & synchronization — Semaphores

#### ➤ Binary semaphore implementation

```
struct binary_semaphore {
    enum {zero, one} value;
    queueType queue;
};

void semWaitB(binary_semaphore s)
{
    if (s.value == one)
        s.value = zero;
    else {
        /* place this process in s.queue */;
        /* block this process */;
    }
}

void semSignalB(semaphore s)
{
    if (s.queue is empty())
        s.value = one;
    else {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```

## c. Concurrency

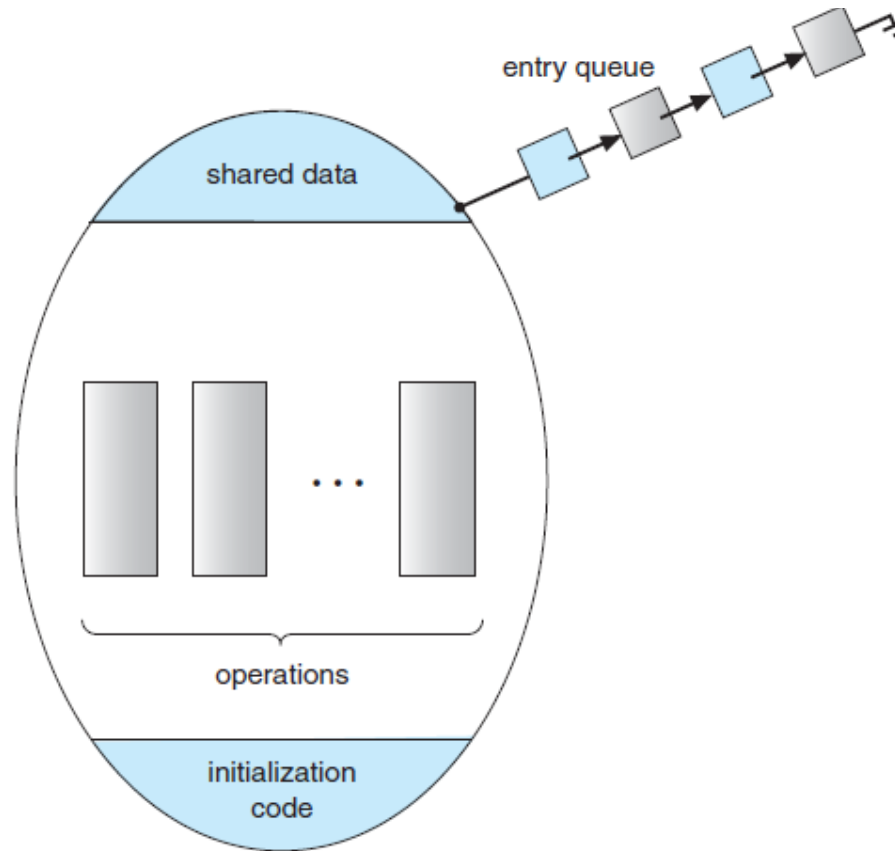
### Mutual exclusion & synchronization — Monitors

- Programming language construct that provides equivalent functionality to that of semaphores and is easier to control
- Implemented in a number of programming languages
  - including Concurrent Pascal, Pascal-Plus, and Java
- Monitors **encapsulate data structures** that are not externally accessible
- **Mutually exclusive** access to data structure enforced by only **“admitting”** one process at a time
  - Processes wanting access join a waiting queue

## c. Concurrency

### Mutual exclusion & synchronization — Monitors

#### ➤ Schematic view of a monitor



- a. Processes (Summary)
- b. Threads (Summary)
- c. Concurrency
  - ✓ Types of process interaction
  - ✓ Race conditions & critical sections
  - ✓ Mutual exclusion by busy waiting
  - ✓ Mutual exclusion & synchronization
- d. **Deadlock**

## d. Deadlock Conditions

- A deadlock situation can arise if the following **4 conditions hold simultaneously**:
  1. Mutual exclusion
  2. Hold and wait
    - a process must be holding **at least one resource** and waiting to acquire additional resources that are currently being held by other processes.
  3. No preemption
    - Resources **cannot be preempted**; that is, a resource can be released only **voluntarily** by the process holding it, after that process has completed its task
  4. Circular wait (Potential consequence of the first three)

## d. Deadlock

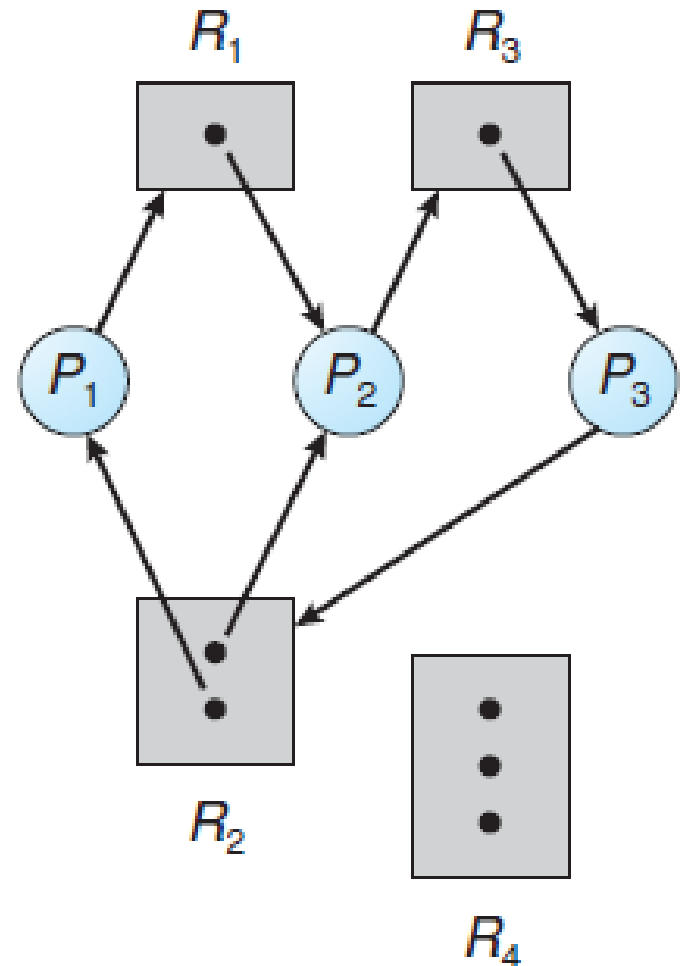
### Resource Allocation Graph (RAG)

**P**: process

**R**: resource

- Instance of a resource

Circular wait:  $P_1$  is waiting for  $R_1$  held by  $P_2$ ,  $P_2$  is waiting for  $R_3$  held by  $P_3$ , and  $P_3$  is waiting for either  $P_2$  or  $P_1$  to release  $R_2$



## d. Deadlock Example

- **Thread A:**  
lock X  
lock Y  
*perform action involving X and Y*  
release X  
release Y
- **Thread B:**  
lock Y  
lock X  
*perform action involving X and Y*  
release X  
release Y

- Possible execution:  
**A: lock X**  
**B: lock Y**  
**A: tries to lock Y**  
**B: tries to lock X**

***DEADLOCK***



## d. Deadlock

### Possible solutions

- Prevent Deadlock
  - adopt a policy that will ensure a deadlock never arises
- Avoid Deadlock
  - avoid entering a state that could lead to deadlock
- Detect Deadlock and Recover

## d. Deadlock

### Prevention approaches

- Design a system in such a way that the possibility of deadlock is **excluded**
- **Direct** or **indirect** approaches

## d. Deadlock

### Indirect Approach to Deadlock Prevention

- Cannot prevent mutual exclusion
- Hold and wait: all resources required to complete must be requested at once; if any one is not available, release all locks
- No pre-emption: if a process holding a resource that is required by another process, OS may pre-empt that resource
- Circular wait: define a linear ordering of resources

## d. Deadlock

### **Direct Approach to Deadlock Prevention**

- Check that any resource request **does not cause a circular wait**

## d. Deadlock

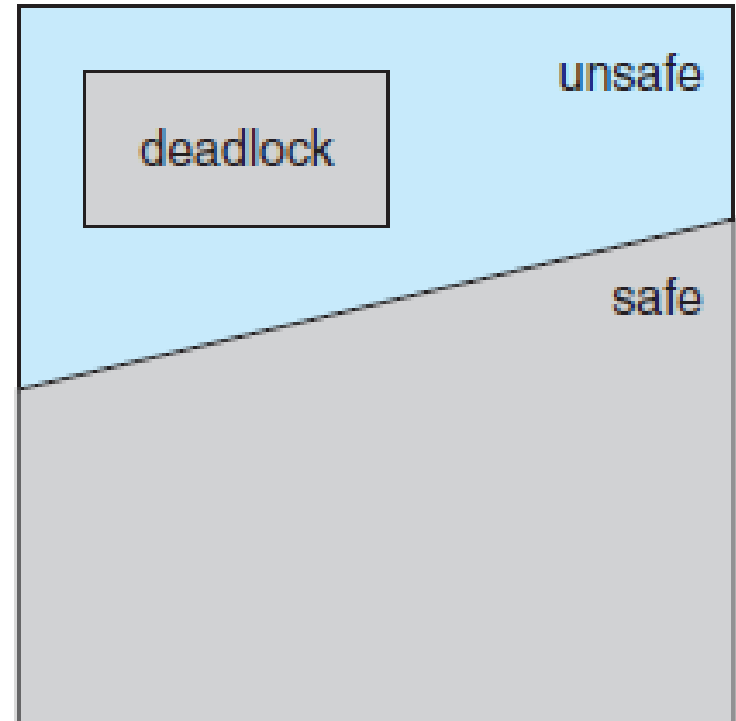
### Deadlock Avoidance

- Process initiation **denial**
  - Do not start a process if **its demands** might lead to deadlock.
- Resource allocation **denial**
  - Do not grant an **incremental resource request** to a process if this allocation might lead to deadlock
  - Example: [banker's algorithm](#)

## d. Deadlock

### Deadlock Avoidance: **Banker's Algorithm**

- **State** of the system reflects the current **allocation of resources to processes**
- **Safe state** is one in which there is **at least one sequence** of resource allocations to processes that **does not result in a deadlock**
- **Unsafe state** is a state that is not safe



## d. Deadlock

### Pros and Cons of Deadlock Avoidance

- It is not necessary to preempt and rollback processes, as in deadlock detection
- It is **less restrictive** than deadlock prevention
- **Max. resource** requirement for each process must be stated in advance
- Processes must be independent; no synch. requirements
- There must be a **fixed number** of resources to allocate
- No process may **exit** while holding resources

## d. Deadlock

### Deadlock **Detection**

- Deadlock detection strategies **do not limit** resource access **or restrict** process actions
- A check for deadlock can be made **as frequently as each resource request** or, less frequently, depending on how likely it is for a deadlock to occur
- The algorithm is relatively simple (Banker's algorithm)
- Frequent checks consume **considerable processor time**



## d. Deadlock

### Deadlock **Recovery**: possible strategies

- **Abort** all deadlocked processes
- **Back up** each deadlocked process to some **previously defined checkpoint** and restart all processes
- **Successively abort** deadlocked processes until deadlock no longer exists
- **Successively preempt resources** until deadlock no longer exists

---

# Next lecture

- Cloud computing