

# Advanced Operating Systems

## Lecture 5: Virtualisation

Soufiene Djahel

Office: John Dalton E151

Email: [s.djahel@mmu.ac.uk](mailto:s.djahel@mmu.ac.uk)

Telephone: 0161 247 1522

Office hours: Monday 10 -11, Thursday 11-13

This lecture was adapted and extended from the slides originally designed by Emma Norling, MMU, UK

# Recap from last week

- Distributed file systems (DFS)
- DFS requirements
- DFS examples:
  - NFS, AFS, NTFS
  - **GFS, HDFS**

# Today's objectives

- To explore the use of **virtualisation** technologies in operating systems
  - Considering **different reasons** for virtualisation
- To describe the **various approaches** to virtualisation
- To discuss **performance issues** in virtualisation

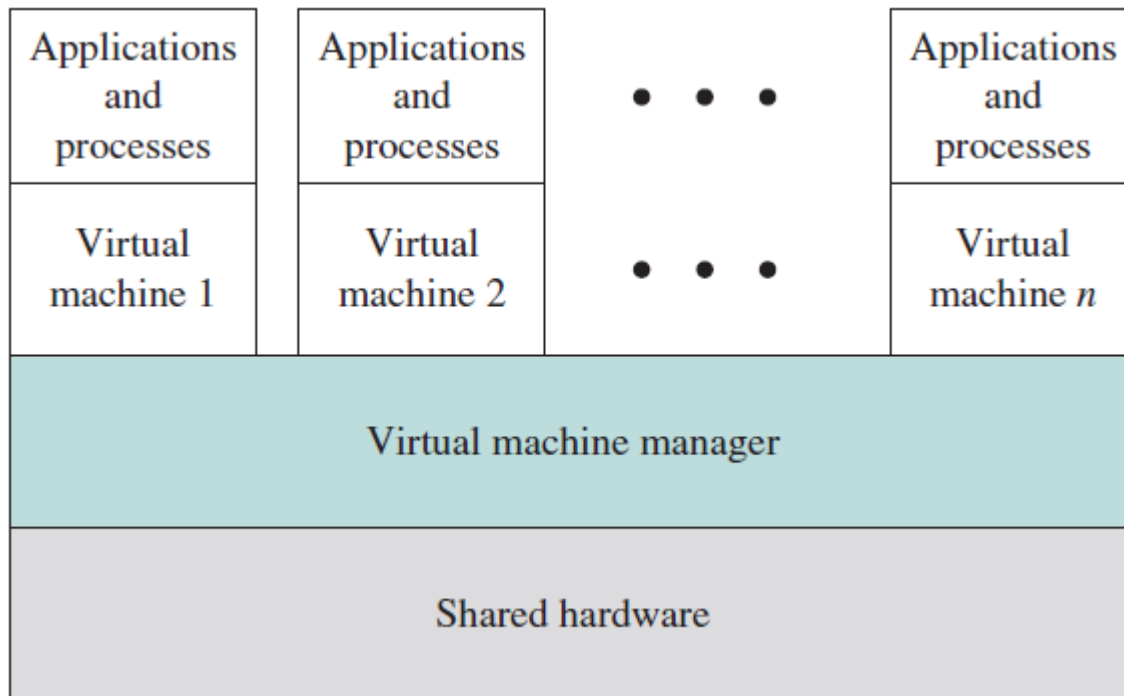
# Virtualisation

- Virtualization technology enables a single PC or server to **simultaneously run multiple** operating systems or multiple sessions of a single OS
- A machine with virtualization software can **host numerous** applications, including those that run on different operating systems, on a **single platform**

## Virtualisation (cont.)

- The solution that enables virtualization is a ***Virtual Machine Monitor (VMM)***, or ***hypervisor***
- This software sits between the hardware and the VMs
  - acting as a **resource broker**
  - allowing multiple VMs to **safely coexist** on a single physical server host and share that host's resources

## Virtualisation (cont.)




**Virtual machine concept**

## Virtualisation (cont.)

- Several components
  - **Host** – underlying hardware system
  - **Virtual Machine Monitor (VMM)** or **hypervisor** – creates and runs virtual machines by providing interface that is *identical* to the host
    - Except in the case of **paravirtualization**
  - **Guest** – process provided with **virtual copy** of the host
    - Usually an operating system

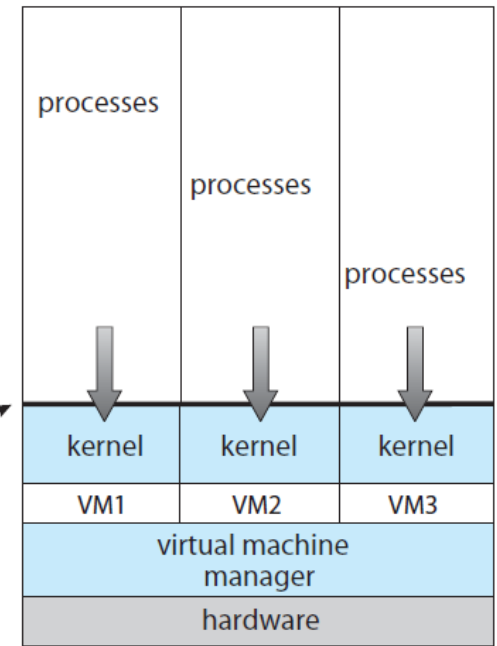
## Virtualisation (cont.)

- Single physical machine can run **multiple** operating systems **concurrently**, each in its **own** **virtual** machine



(a)

Non-virtual machine




(b)

Virtual machine



# Consolidation ratio

- The number of Guests that can exist on a single host is measured as **Consolidation ratio**
  - e.g., a host supporting 6 VMs has a consolidation ratio of 6 to 1, written as **6:1**



## Virtual machine consolidation

# History

- Virtual machines invented by IBM in late 1960s
- Original usage:
  - Each user ran a different **guest OS**
  - Single **shared hardware platform**
- Interest died out in the 1980's and 1990's:
  - Each user had a private machine
- Reinvented, made practical by Mendel Rosenblum and graduate students at Stanford, led to **VMware**.

# Why Use Virtualisation?

- Cloud computing (**laaS**)
- Multiple OSes on one machine
  - **Multi-platform** development/support
- Fault **isolation**
- Encapsulation

# Encapsulation

- VMM can encapsulate **entire state** of a VM in a file.
- Can **save, continue, restore** old state.
- Datacenter example:
  - Can **migrate** VM's between machines to **balance load**
- Software development:
  - Tests may **corrupt** the state of the machine
  - Solution:
    - Run tests in a VM
    - Always start tests from a **saved VM configuration**
    - Discard VM state after tests
    - Results: **reproducible tests**

# Supporting Virtualisation

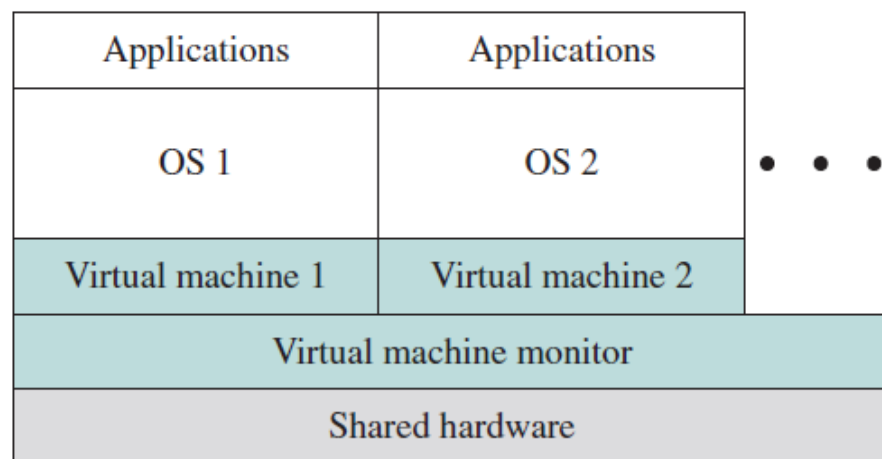
- Need a process that is an **abstraction** of a machine, which looks exactly like the underlying hardware:
  - The complete **instruction set** of the underlying machine
  - Physical memory
  - Memory management unit (page tables, etc.)
  - I/O devices
  - Traps and interrupts
- **Virtual Machine Monitor** (VMM) (aka *hypervisor*)

# Virtual Machine Monitor (VMM)

- Implements **sharing of real hardware resources** by multiple OS VMs that each think they have a complete copy of the machine
  - Popular in early days of computing to allow mainframe to be shared by multiple groups developing OS code
  - Used in modern mainframes to allow multiple versions of OS to be **running simultaneously** → OS upgrades with no downtime!
  - Example for PCs: **VMware** allows Windows OS to run on top of Linux (or vice-versa)

# VMM types

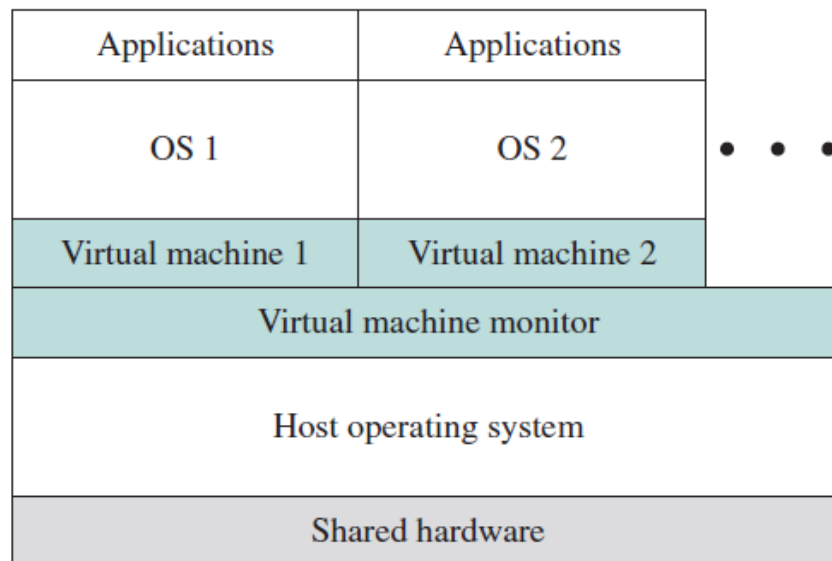
- Type 1 VMM, aka “bare metal” VMM
  - Loaded as a thin software layer into the physical server
  - VMM is the **host OS**
  - Thus, it can **directly control** the resources of the host
  - E.g., VMware ESXi, Microsoft Hyper-V



(a) Type 1 VMM

# VMM types

- Type 2 VMM, hosted by the underlying OS
  - Works as a **traditional app**, a program code loaded on top of OS
  - Relies on the OS to handle all hardware interactions
  - It has **lower performance** than Type 1 VMM
  - E.g., VMware Workstation, ORACLE Virtual Box



(b) Type 2 VMM




# Approaches to VMMs

- **Simulation**
  - sometimes called **emulation**, but it is not full emulation
- Use CPU, but **trap** on access to privileged hardware state
- Paravirtualisation
- Hardware-assisted virtualisation

# Evolution of Software solutions\*

- **1<sup>st</sup> Generation:** Full virtualization (**Binary rewriting**)
  - Software Based
  - VMware and Microsoft
- **2<sup>nd</sup> Generation:** Paravirtualization
  - **Cooperative** virtualization
  - **Modified** guest
  - VMware, Xen
- **3<sup>rd</sup> Generation:** Silicon-based (Hardware-assisted) virtualization
  - **Unmodified** guest
  - VMware and Xen on virtualization-aware hardware platforms



\*This slide is from Intel® Corporation

# Simulation

- Most **basic** form of virtualisation
- All underlying hardware is simulated
  - e.g. QEMU
- SLOW

# Instruction Traps


- **Most common** form of virtualisation
- Guest OS runs as an application program
  - Attempts to access privileged instructions cause exceptions
- When guest OS issues a privileged instruction, **VMM traps it**, takes appropriate action, and returns result

# Paravirtualisation

- Used to **improve performance**
- Guest OSes “know” that they are **modified**
- Make use of specialist instructions to improve **speed** and **reliability**
- Must be supported by underlying hardware *and* guest OS

# Hardware-assisted Virtualization

- Server hardware is virtualization aware
- VMM loads at privilege Ring -1 (**firmware**)
- Removes CPU **emulation bottleneck**
- Memory virtualization coming in quad core AMD and Intel CPUs



Hardware-assisted virtualization

# What If We Can't Trap...?

- If there's no way to tell when guest OS is trying to access **privileged instructions**...
  - One possibility: VMM checks *every* instruction when it is executed
    - **costly**, but not as bad as simulation
  - Binary translation: Instruction set translated in advance to replace sections that require virtualisation
    - **Upfront cost**, but can also be optimised

- Although VMM isolates VMs, they can affect each other
  - resource usage
- Different VMMs manage resource usage in different ways
- In cloud computing, additional options of *migration exist*
  - VM moves from one **VMM** (and one **underlying hardware platform**) to another



# VMM overhead

- Depends on **workload**
- Goal for system VMs:
  - Run almost all instructions directly on **native hardware**
- User-level CPU-bound programs have **near-zero** virtualization overhead
  - Run at **native speeds** since OS rarely invoked

## VMM overhead (cont.)

- I/O-intensive workloads are **OS-intensive**
  - Execute many **system calls** and **privileged instructions**
  - Can result in **high virtualization overhead**
- But if I/O-intensive workload is also **I/O-bound**
  - Processor utilization is low (since waiting for I/O)
  - Processor virtualization can be **hidden** in I/O costs
  - So virtualization overhead is **low**

---

# Full Virtualisation

# Processor Management

- Two main strategies:
  1. Emulate chip as software (e.g. QEMU or Android Emulator in Android SDK).
    - Easily transportable but inefficient
  2. Provide each VM with segments of time on the physical processors (**pCPUs**) to host the virtual processors (**vCPUs**)
    - Number of processors (**vCPUs**) is important metric in system
    - In general, CPUs are **underutilised** → start with one, add extras if needed

# Processor Management (cont.)

- Operating systems in general have at least **two levels** of protection: **kernel mode** and **user mode**
- Most systems have **four levels** of protection, described as rings:
  - **Ring 0** is most privileged, where kernel runs
  - **Ring 3** is least privileged, where user processes run
  - **Rings 1 & 2** in theory for intermediate levels of protection, such as device drivers
- A guest OS believes to be operating at **ring 0**, but in reality is only doing so on the virtualised machine
  - In terms of the underlying hardware, it's operating in **ring 3**

# Memory Management

- A VM needs to be configured with enough memory to function efficiently
- If each guest OS in every VM **manages its own set of page tables**, what happens to the virtual memory managed by the underlying OS?

# Memory Management (cont.)

- Three abstractions of memory:
  - **Real (aka machine) memory:** actual hardware memory, e.g., 8GB of DRAM
  - **Physical memory:** abstraction of hardware memory managed by a guest OS
    - If the VMM allocates **512 MB** to a VM, its OS thinks the computer has **512MB** of contiguous physical memory
  - **Virtual memory:**
    - a continuous **virtual address space** presented by the guest OS to applications running inside the VM.

# Memory Management (cont.)

- VMM separates **real** and **physical** memory
  - Makes **real** memory a separate, intermediate level between **virtual** memory and **physical** memory
  - Guest OS maps virtual memory to physical memory via its **page tables**
  - **VMM page tables** map real memory to physical memory



## Memory Management (cont.)

- VMM manages **page sharing**, so the virtual machine OSs are **unaware** of what is happening in the physical system
  - Multiple VMs share pages if they run similar version of OS or the same applications

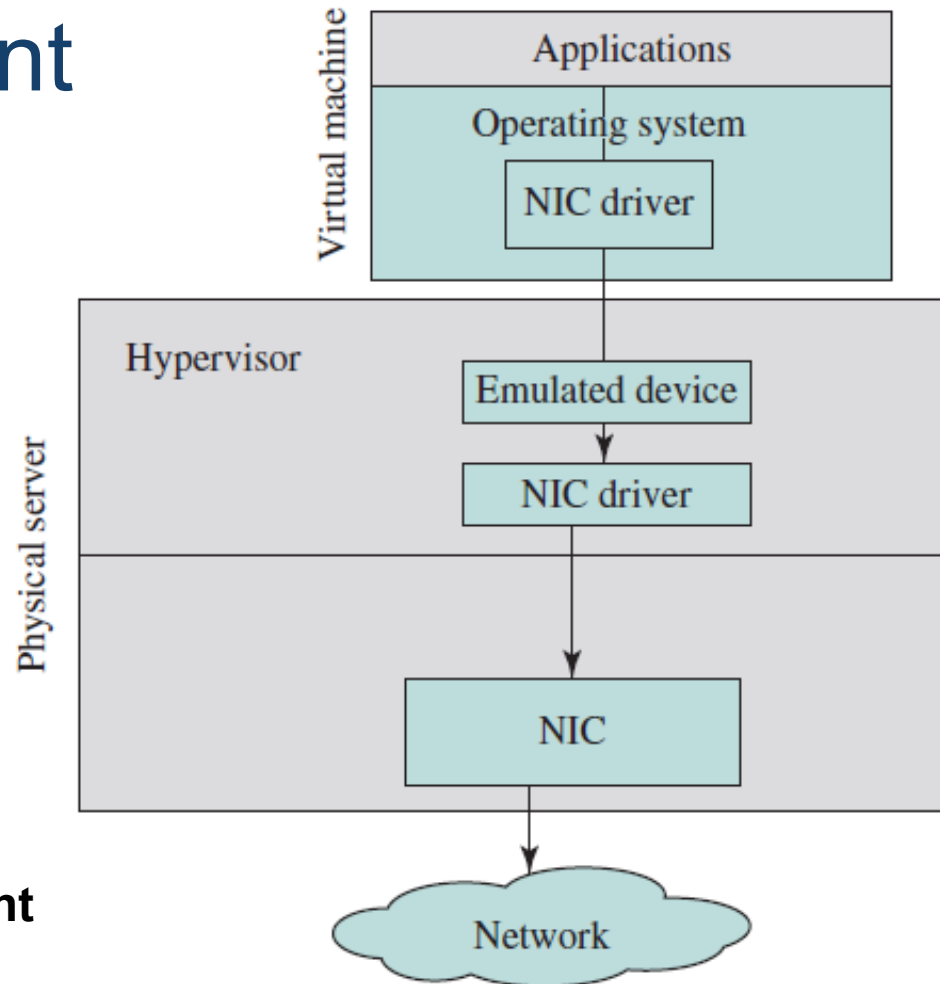
# Memory Management (cont.)

- Ballooning
  - the VMM activates a balloon driver that (virtually) **inflates** and presses the guest OS to flush pages to disk
  - once the pages are cleared, the balloon driver **deflates** and the VMM can use the physical memory for other VMs
- Memory **overcommit**
  - the capability to allocate more memory than physically exists on a host

# I/O Management

- **Most difficult** part of virtualization
  - Increasing number of I/O devices attached to the computer
  - Increasing diversity of I/O device types
  - Sharing of a real device among multiple VMs,
  - Supporting the myriad of device drivers that are required, especially if different guest OSes are supported on the same VM system

# I/O Management



**I/O in a virtual environment**

# I/O Management

- An advantage of virtualizing the workload's I/O path is enabling **hardware independence** by abstracting vendor-specific drivers to more generalized versions that run on the hypervisor
- This abstraction enables:
  - **live migration**, which is one of virtualization's greatest availability strengths
  - the sharing of **aggregate resources**, such as network paths

# I/O Management

- The memory **overcommit** capability is another benefit of virtualizing the I/O of a VM
- The trade-off for this is that the hypervisor is managing all the traffic **and requires processor overhead**
  - this was an issue in the early days of virtualization but now faster multicore processors and sophisticated hypervisors have addressed this concern

---

# VMM Example

# Java VM – not a VM in the same sense!

- The goal of a Java Virtual Machine (JVM) is to provide a **runtime space** for a set of Java code to run on **any operating system** staged on **any hardware platform** without needing to make code changes to accommodate the different operating systems or hardware
- The JVM can support multiple threads
- Promises “**Write Once, Run Anywhere**”



# Java VM

- Example of programming-environment virtualization
- Includes language specification (Java), API library, **Java virtual machine (JVM)**
- Java objects specified by **class construct**, Java program is one or more objects

# Java VM

- The JVM is described as being an **abstract computing machine** consisting of:
  - an instruction set
  - a program counter register
  - a stack to hold variables and results
  - a heap for runtime data
  - a method area for code and constants

# Android Virtual Machine

- Referred to as Dalvik
- The **Dalvik VM** (DVM) executes files in the Dalvik Executable (.dex) format
- The Dalvik core class library is intended to provide a familiar development base for those used to programming with Java Standard Edition, but it is geared specifically to the needs of a small mobile device
- Each Android application runs in its own process, with its **own instance** of the Dalvik VM
- Dalvik has been written so that a device can run multiple VMs efficiently

# Summary

- Range of reasons for virtualisation
  - Cloud computing is just one
- Various approaches to virtualisation
- Virtualisation leads to performance issues
- Range of commercial and open-source examples

# Interesting VM tools

- Xen:
  - <https://www.xenproject.org/>
  - <https://xenserver.org/>
- VMWare: <https://www.vmware.com/>
- VirtualBox: <https://www.virtualbox.org/>

---

# Next lecture: Fault tolerance